

Chap 7 Ulysses' Compass

```
1 #     margin-left: 1%;  
2 #     margin-right: 5%;  
3 html"""  
4 main {  
5     margin: 0 auto;  
6     max-width: 90%;  
7     padding-left: max(50px, 1%);  
8     padding-right: max(253px, 10%);  
9     # 253px to accomodate TableOfContents(aside=true)  
10 }  
11 """"
```

```
1 versioninfo()
```

```
Julia Version 1.10.2  
Commit bd47eca2c8a (2024-03-01 10:14 UTC)  
Build Info:  
    Official https://julialang.org/ release  
Platform Info:  
    OS: Linux (x86_64-linux-gnu)  
    CPU: 32 × Intel(R) Xeon(R) CPU E5-2630 v3 @ 2.40GHz  
    WORD_SIZE: 64  
    LIBM: libopenlibm  
    LLVM: libLLVM-15.0.7 (ORCJIT, haswell)  
Threads: 16 default, 0 interactive, 8 GC (on 32 virtual cores)  
Environment:  
    JULIA_PKG_SERVER = https://mirrors.tuna.tsinghua.edu.cn/julia  
    JULIA_REVISE_WORKER_ONLY = 1
```

```
1 using Pkg, DrWatson
```

Table of Contents

Chap 7 Ulysses' Compass

link function from StatisticalRethinking.

link

Required arguments

Return values

Examples

7.0 How to display stdout/terminal output in Pluto.jl

7.1 The problem with parameters.

7.1 Load data & 7.2 Rescale data

Code 7.3 Regress brain on mass, optimized via MAP

Checking the m7_1 Chains

Code 7.4 OLS to obtain posterior distribution of a & b, but not σ

Code 7.5 calculate R^2

Do explicit simulation due to $\log \sigma$.

Code 7.6 Why R^2 is bad?

Code 7.7 2nd-degree polynomial

Code 7.8 3-, 4-, 5-degree polynomial fitting: $\text{brain_std} \sim \text{mass_std}$

Code 7.9: 6-degree polynomial

Code 7.10 Fig7.3

Selection deleted

Reimplemented the brain_plot function to check my results.

Code 7.11

7.2 Entropy and accuracy.

Code 7.12 Information and KL Divergence

7.13 lppd: Log-Pointwise-Predictive-Density

7.14 Run lppd manually

7.15 lppd function

7.16 Define multiple functions

7.17 Multi-threading

7.18 Plot the training and testing deviance

7.3 Golem taming: regularization

7.4 Predicting predictive accuracy

7.19 Monte Carlo of a Bayesian Linear Model

7.20 logprob

7.21 lppd

7.22 penalty of WAIC

7.23 WAIC

7.24 stddev of WAIC

7.5 Model comparison

7.25 WAIC for m6_7 (include both treatment and fungus)

7.26 Compare WAIC of m6.6 - m6.8

7.27 stddev of WAIC delta between m6.7 and m6.8

7.28 99% confidence interval of dWAIC

7.29 plot WAIC/deviance of 3 different models

7.30 stddev of WAIC differences between m6.6 and m6.8

7.31 stddev of WAIC differences among all models

7.32 Fit 3 models of the divorce dataset (Divorce rate ~ Marriage rate, ...)

7.33 Compare 3 models via PSIS: m5_1 is best.

7.34 m5_3: Compare pointwise PSIS Pareto k and WAIC penalty

7.34.1 m5_3_log_sigma_normal: Compare pointwise PSIS Pareto k and WAIC...

7.35 m5_3t: MV T-dist df=2, Compare pointwise PSIS Pareto k and WAIC...

7.36 m5_3t: 1D T-dist (failed in ForwardDiff but succeeded via using ar...

Try Truncated cauchy as σ (represented as $\log \sigma$ so that estsummary ca...

7.37 m5_1t: m5_1 but uses T-distribution

```
1 begin
2   using PlutoUI
3   TableOfContents()
```

```
1 begin
2     using Optim
3     using GLM
4     using CSV
5     using Random
6     using StatsBase
7     using DataFrames
8     using Dagitty
9     using Turing
10    using StatsPlots
11    using StatisticalRethinking
12    using StatisticalRethinkingPlots
13    using ParetoSmoothedImportanceSampling
14    import ParetoSmoothedImportanceSampling as psis
15    import ParetoSmooth
16    using Logging
17    # DocStringExtensions provides $(SIGNATURES)
18    using DocStringExtensions
19    #using ProgressBars
20    # have to import the student T distribution explicitly,
21    # as it is not exported
22    import Distributions: IsoTDist
end
```

Selection deleted

link function from StatisticalRethinking.

- The variant that takes a function as input.

link

link

Generalized link function to evaluate callable for all parameters in dataframe over range of x values.

```
link(dfa, rx_to_val, xrange)
```

Required arguments

- `dfa::DataFrame`: data frame with parameters
- `rx_to_val::Function`: function of two arguments: row object and x
- `xrange`: sequence of x values to be evaluated on

Return values

Selection deleted

Is the vector, where each entry was calculated on each value from `xrange`. Every such entry is a list corresponding each row in the data frame.

Examples

```
julia> using StatisticalRethinking, DataFrames
julia> d = DataFrame(:a => [1,2], :b=>[1,1])
2×2 DataFrame
 Row | a      b
     | Int64  Int64
   ──
  1 |      1      1
  2 |      2      1
julia> link(d, (r,x) -> r.a+x*r.b, 1:2)
2-element Vector{Vector{Int64}}:
 [2, 3]
 [3, 4]
```

```
1 """
2 # link
3
4 Generalized link function to evaluate callable for all
5 parameters in dataframe over range of x values.
6
7 $(SIGNATURES)
8
9 ## Required arguments
10 * `dfa::DataFrame`: data frame with parameters
11 * `rx_to_val::Function`: function of two arguments: row
12 object and x
13 * `xrange`: sequence of x values to be evaluated on
14
15 ## Return values
16 Is the vector, where each entry was calculated on each
17 value from xrange.
18 Every such entry is a list corresponding each row in the
19 data frame.
20
21 ## Examples
22 """jldoctest
```

```

22 julia> using StatisticalRethinking, DataFrames
23
24 julia> d = DataFrame(:a => [1,2], :b=>[1,1])
25 2×2 DataFrame
26 Row | a      b
27     | Int64  Int64
28
29 1 | 1      1
30 2 | 2      1
31
32 julia> link(d, (r,x) -> r.a+x*r.b, 1:2)
33 2-element Vector{Vector{Int64}}:
34 [2, 3]
35 [3, 4]
36
37 """
38 """
39 function link(dfa::DataFrame, rx_to_val::Function, xrange)
40 [
41   rx_to_val.(eachrow(dfa), (x,))
42   for x ∈ xrange
43 ]

```

7.0 How to display stdout/terminal output in Pluto.jl

```

1 begin
2   Plots.default(labels=false)
3   # Disable the following line first.
4   #Logging.disable_logging(Logging.Warn);
5 end;

```

- The following for-loop displays nothing because the loop returns a nothing object.
- Pluto.jl only displays returned object.

```

1 for i in 1:20
2   R = i+100
3   T = R^2
4   md"hello $R and $T"
5 end

```

- the following for-loop will return and output the last iteration. but you can run it differently via map(...) do ... end

hello 120 and 14400

```

1 let
2   x = nothing
3   for i in 1:20
4     R = i+100
5     T = R^2
6     # something
7     x = md"hello $R and $T"
8   end
9   x
10 end

```

- Output of every iteration of the for-loop is possible due to map(1:20)

```
▶ [hello 101 and 10201, hello 102 and 10404, hello 103 and 10609,
```

```
1 map(1:20) do i
2   R = i+100
3   T = R^2
4   md"hello $R and $T"
5 end
```

```
@with_stdout (macro with 1 method)
```

```
1 macro with_stdout(expr)
2   escaped_expr = esc(expr)
3   return quote
4     stdout_bk = stdout
5     rd, wr = redirect_stdout()
6     result = ($escaped_expr)
7     redirect_stdout(stdout_bk)
8     close(wr)
9     print_result = read(rd, String) |> Text
10    (result, print_result)
11  end
12 end
```

```
1 @with_stdout println("I am");
```

~~Selection deleted~~ String (enable Logging.Warn), terminal output is now displayed in Pluto.jl

```
1 @time println("I am");
```

```
I am
0.000103 seconds (18 allocations: 200.500 KiB)
```

```
@seeprints (macro with 1 method)
```

```
1 macro seeprints(expr)
2   quote
3     stdout_bk = stdout
4     rd, wr = redirect_stdout()
5     $expr
6     redirect_stdout(stdout_bk)
7     close(wr)
8     read(rd, String) |> Text
9   end
10 end
```

```
1 @seeprints println("I am");
```

7.1 The problem with parameters.

7.1 Load data & 7.2 Rescale data

```
1 begin
2   sppnames = ["afarensis", "africanus", "habilis",
3   "boisei",
4   "rudolfensis", "ergaster", "sapiens"]
5   brainvolcc = [438, 452, 612, 521, 752, 871, 1350]
6   masskg = [37.0, 35.5, 34.5, 41.5, 55.5, 61.0, 53.5]
7   d = DataFrame(:species => sppnames, :brain =>
8   brainvolcc, :mass => masskg)
9   d[!, :mass_std] = (d.mass .- mean(d.mass)) ./ std(d.mass)
10  d[!, :brain_std] = d.brain ./ maximum(d.brain)
11  d
12 end;
```

Code 7.3 Regress brain on mass, optimized via MAP

To match results from the book, the model was optimized using MAP estimation, not the MCMC I used before. The reason for that is the MCMC producing different estimation for \log_{σ} value, which makes all the values very different.

If you want, you can experiment with NUTS sampler for models in 7.1 and 7.2.

In addition, you can also check [this discussion](#)

```
model_m7_1 (generic function with 2 methods)
1 @model function model_m7_1(mass_std, brain_std)
2   a ~ Normal(0.5, 1)
3   b ~ Normal(0, 10)
4   μ = @. a + b*mass_std
5   log_σ ~ Normal()
6   brain_std ~ MvNormal(μ, exp(log_σ))
7 end
```

Selection deleted

	variable	mean	min	median	max	nmissi
1	:a	0.53284	-0.000688679	0.534284	1.07771	0
2	:b	0.170031	-0.174648	0.16508	0.607495	0
3	:log_σ	-1.37747	-2.37392	-1.40897	0.42064	0

```
1 begin
2   @time m7_1_ch = sample(model_m7_1(d.mass_std,
3     d.brain_std), NUTS(), 1000)
4   m7_1 = DataFrame(m7_1_ch)
5   @time describe(m7_1)
6 end
```

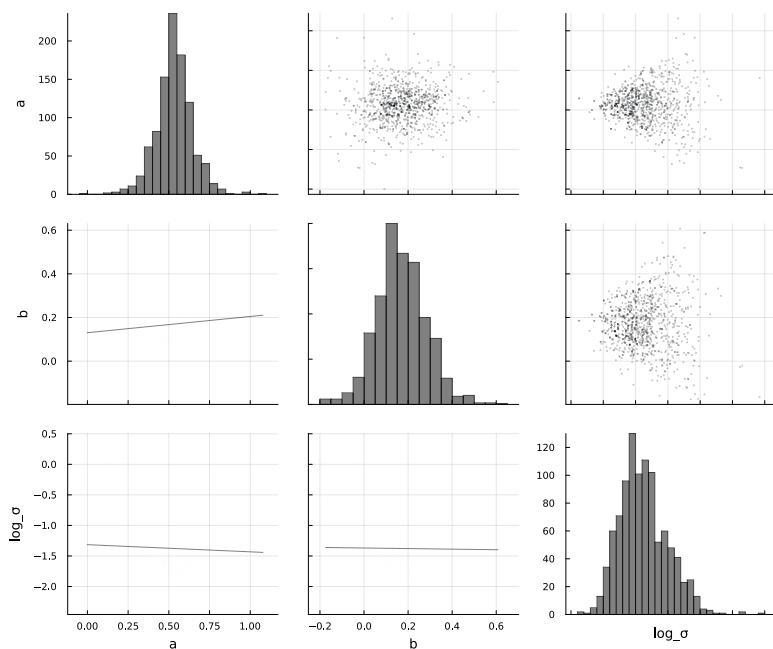
100%

ⓘ Found initial step size
ε: 0.4

4.247740 seconds (3.32 M allocations: 249.010 MiB, 2.40% gc time, 68.20% compilation time)
0.000213 seconds (118 allocations: 30.766 KiB)

▶ (1000, 3)

```
1 size(m7_1)
```



```
1 @df DataFrame(m7_1) corrplot(cols(1:3),
  series_type=:scatter, ms=0.2, alpha=0.5, size=(950, 800),
  bins=30, grid=true)
```

Selection deleted

Checking the m7_1 Chains

```
1 md"### Checking the m7_1 Chains"
```

	iteration	chain	a	b	log_sigma	lp	l
1	501	1	0.509843	0.0890367	-2.04102	-5.61041	:
2	502	1	0.509843	0.0890367	-2.04102	-5.61041	:
3	503	1	0.555917	0.145738	-1.76824	-3.80884	:
4	504	1	0.555917	0.145738	-1.76824	-3.80884	:
5	505	1	0.530114	0.146504	-1.61682	-3.72554	:
6	506	1	0.775115	0.0608116	-1.60618	-9.86268	:
7	507	1	0.738181	0.189116	-1.27703	-6.51151	:
8	508	1	0.651913	0.117101	-0.36206	-9.33962	:
9	509	1	0.686841	0.36336	-1.38415	-7.40292	:
10	510	1	0.686841	0.36336	-1.38415	-7.40292	:
	:: more						

```
1 m7_1_ch
```

```
Chains{Float64, AxisArrays.AxisArray{Float64, 3, Array{Float64, 3}}}
```

```
1 typeof(m7_1_ch)
```

```
▶ (varname_to_symbol = OrderedDict(a => :a, b => :b, log_sigma => :log
```

```
1 m7_1_ch.info
```

```

3-dimensional AxisArray{Float64,3,...} with axes:
  :iter, 501:1:1500
  :var, [:a, :b, :log_σ, :lp, :n_steps, :is_accept, :acceptance
  :chain, 1:1
And data, a 1000×15×1 Array{Float64, 3}:
[:, :, 1] =
0.509843  0.0890367 -2.04102 ...
0.509843  0.0890367 -2.04102  4.27759  2.0  0.0  0.64834
0.555917  0.145738  -1.76824  2.75461  2.0  0.0  0.64834
0.555917  0.145738  -1.76824  0.677042  1.0  0.0  0.64834
0.530114  0.146504  -1.61682  0.53324  2.0  0.0  0.64834
0.775115  0.0608116 -1.60618 ...
0.738181  0.189116  -1.27703 -2.59115  2.0  0.0  0.64834
...
0.466506  0.0977929 -1.84985 -0.19042  2.0  0.0  0.64834
0.524854  0.205494  -2.05463 -0.685438 2.0  0.0  0.64834
0.524854  0.205494  -2.05463  4.72687  2.0  0.0  0.64834
0.471906  0.225656  -1.67074  6.68196  2.0  0.0  0.64834
0.518771  0.0432336 -1.29095  0.387717 2.0  0.0  0.64834
0.491565  -0.00268598 -1.40619  0.211685  1.0  0.0  0.64834

```

```

1 begin
2   @show m7_1_ch.name_map
3   m7_1_ch.value
4 end

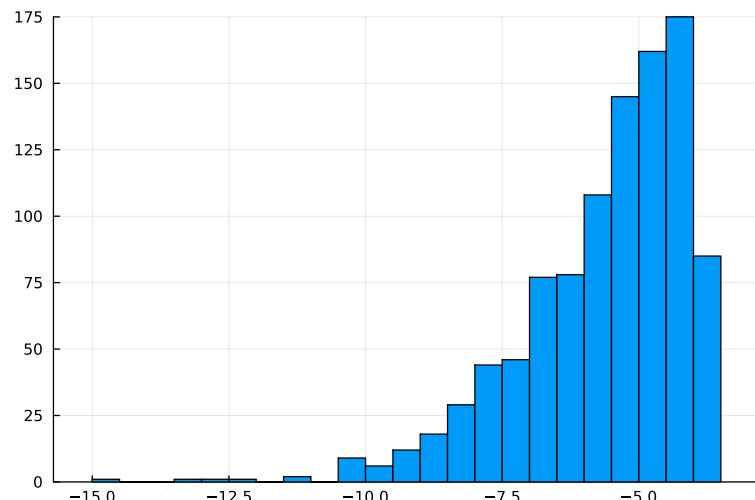
```

Select

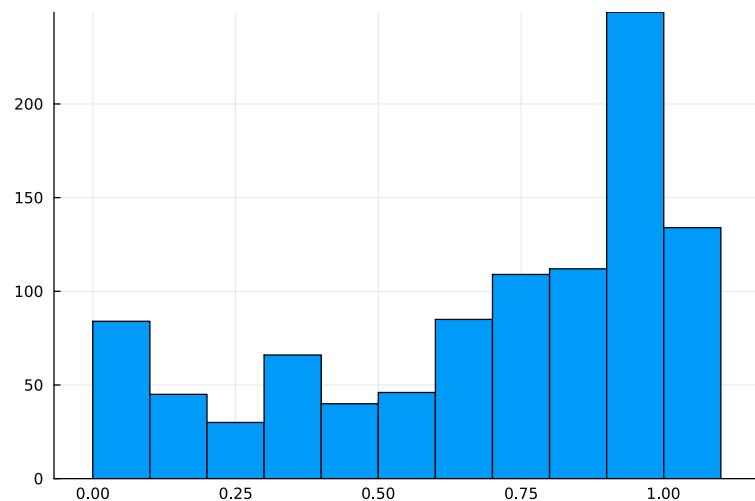
```

m7_1_ch.name_map = (parameters = [:a, :b, :log_σ],
internals = [:lp, :n_steps, :is_accept, :acceptance_rate,
:log_density, :hamiltonian_energy, :hamiltonian_energy_error,
:max_hamiltonian_energy_error, :tree_depth,
:numerical_error, :step_size, :nom_step_size])

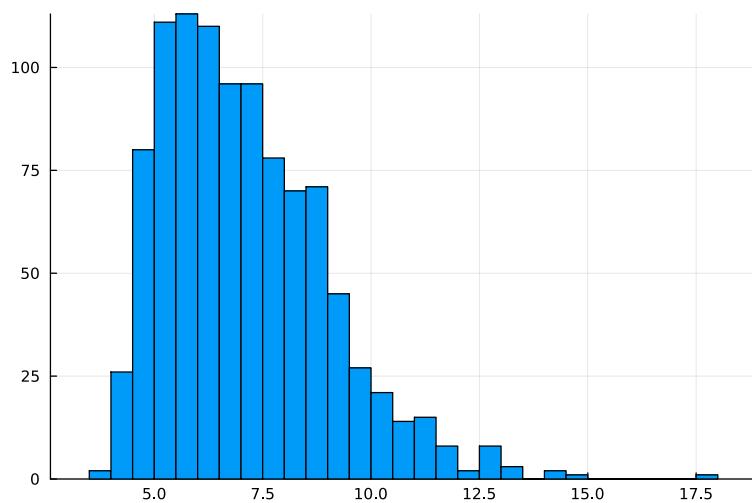
```



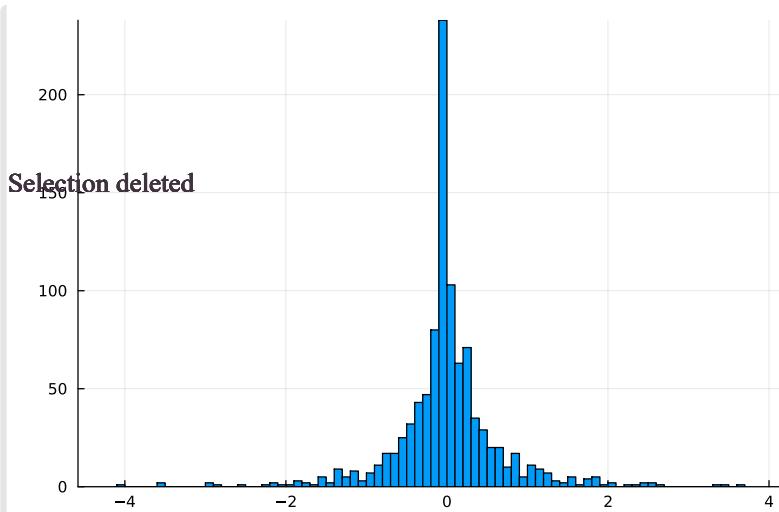
```
1 histogram(m7_1_ch[:log_density])
```



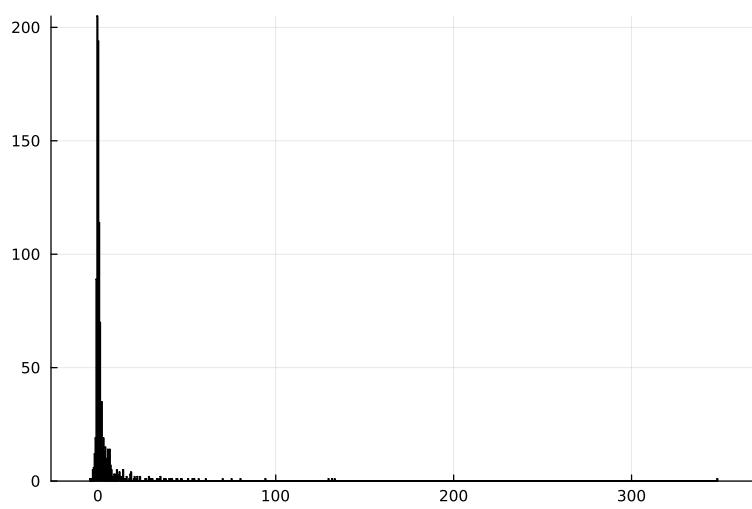
```
1 histogram(m7_1_ch[:acceptance_rate])
```



```
1 histogram(m7_1_ch[:hamiltonian_energy])
```



```
1 histogram(m7_1_ch[:hamiltonian_energy_error])
```



```
1 histogram(m7_1_ch[:max_hamiltonian_energy_error])
```

```

2-dimensional AxisArray{Float64,2,...} with axes:
  :iter, 501:1:1500
  :chain, 1:1
And data, a 1000×1 Matrix{Float64}:
0.6483448249395375
0.6483448249395375
0.6483448249395375
0.6483448249395375
0.6483448249395375
0.6483448249395375
0.6483448249395375
0.6483448249395375
0.6483448249395375
0.6483448249395375
⋮
0.6483448249395375
0.6483448249395375
0.6483448249395375
0.6483448249395375
0.6483448249395375
0.6483448249395375
0.6483448249395375

```

1 m7_1_ch[:nom_step_size]

Code 7.4 OLS to obtain posterior distribution of a & b, but not σ

```

X = 7×2 Matrix{Float64}:
 1.0  -0.779467
 1.0  -0.91702
 1.0  -1.00872
Selection deleted 66808
 1.0  0.91702
 1.0  1.42138
 1.0  0.733616
1 X = hcat(ones(length(d.mass_std)), d.mass_std)

m =
GLM.LinearModel{GLM.LmResp{Vector{Float64}}}, GLM.DensePredChol{Fl
Coefficients:

```

	Coef.	Std. Error	t	Pr(> t)	Lower 95%	Upper 95%
x1	0.528677	0.0705692	7.49	0.0007	0.347273	0.710081
x2	0.167118	0.0762235	2.19	0.0798	-0.0288204	0.363057

1 m = lm(X, d.brain_std)

Code 7.5 calculate R^2

1 Random.seed!(12);

Do explicit simulation due to \log_σ .

- Why?

```

0.48838799760488294
1 @time begin
2     # d.mass_std is a 7-element vector.
3     s = [
4         rand(MvNormal(@. r.a + r.b * d.mass_std),
5             exp(r.log_σ)))
6         for r ∈ eachrow(m7_1)
7     ]
8     # s is a vector of length 1000. each element is a 7-
9     # element vertical vector.
10    @show size(s)
11    # transpose each element of s to horizontal vector and
12    # then vertically concatenate them.
13    s = vcat(s'...);
14    @show size(s)
15
16    r = mean.(eachcol(s)) .- d.brain_std;
17    @show size(r)
18    @show resid_var = var(r, corrected=false)
19    outcome_var = var(d.brain_std, corrected=false)
20    @show outcome_var
21    1 - resid_var/outcome_var
end

```

```

size(s) = (1000,)
size(s) = (1000, 7)
size(r) = (7,)
resid_var = var(r, corrected = false) = 0.024986533810
43236
outcome_var = 0.048838834299151763
0.004239 seconds (18.69 k allocations: 2.092 MiB)

```

```

simulate(df::DataFrames.DataFrame, rx_to_dist::Function, xrange;
return_dist, seed) in StatisticalRethinking at
/y/home/huangyu/julia/packages/StatisticalRethinking/RYYWV/src/simula
1 @which simulate(m7_1, (r, x) -> Normal(r.a + r.b * x,
exp(r.log_σ)), d.mass_std)

```

```

IsoNormal(
dim: 3
μ: [1.0, 2.0, 3.0]
Σ: [0.01000000000000002 0.0 0.0; 0.0 0.01000000000000002 0.0; 0.0 0.0 0.01000000000000002]
)
1 MvNormal([1,2,3], 0.1)

```

Code 7.6 Why R^2 is bad?

Function is implemented in a generic way to support any amount of $b[x]$ coefficients.

```
R2_is_bad (generic function with 1 method)
1 function R2_is_bad(df; sigma=missing)
2     degree = ncol(df[!,r"b"])
3     # build mass_std*degree matrix, with each col
4     # exponentiated to col's index
5     t = repeat(d.mass_std, 1, degree)
6     t = hcat(map(.^, eachcol(t), 1:degree)...)
7     s =
8         begin
9             # calculate product on coefficient's vector
10            b = collect(r[r"b"])
11            μ = r.a .+ t * b
12            s = ismissing(sigma) ? exp(r.log_σ) : sigma
13            rand(MvNormal(μ, s))
14        end
15        for r ∈ eachrow(df)
16    ]
17    s = vcat(s'...);
18 Selection deleted
19 r = mean.(eachcol(s)) .- d.brain_std;
20 v1 = var(r, corrected=false)
21 v2 = var(d.brain_std, corrected=false)
22 1 - v1 / v2
end
```

Code 7.7 2nd-degree polynomial

```
model_m7_2 (generic function with 2 methods)
1 @model function model_m7_2(mass_std, brain_std)
2     a ~ Normal(0.5, 1)
3     b ~ MvNormal([0, 0], 10)
4     μ = @. a + b[1]*mass_std + b[2]*mass_std^2
5     log_σ ~ Normal()
6     brain_std ~ MvNormal(μ, exp(log_σ))
7 end
```

variable	mean	min	median	max
1 :a	0.607718	-1.40189	0.606191	2.37658
2 Symbol("b[1]")	0.194592	-2.13629	0.195842	1.44602
3 Symbol("b[2]")	-0.0968602	-2.14742	-0.0949097	1.66151
4 :log_σ	-1.27554	-2.24776	-1.31453	1.32956

```

1 begin
2   @time m7_2_ch = sample(model_m7_2(d.mass_std,
3     d.brain_std), NUTS(), 10000)
4   @time m7_2 = DataFrame(m7_2_ch)
5   @time describe(m7_2)
end

100%
① Found initial step size
  ε: 0.4

8.853600 seconds (16.16 M allocations: 1.366 GiB, 4.16% gc time, 36.34% compilation time)
Selection 0.0002215 seconds (113 allocations: 631.438 KiB)
  0.000739 seconds (137 allocations: 320.500 KiB)

```

Code 7.8 3-, 4-, 5-degree polynomial

fitting: brain_std ~ mass_std

Implemented the sample in a general way.

```

model_m7_n (generic function with 2 methods)
1 @model function model_m7_n(mass_std, brain_std; degree::Int)
2   a ~ Normal(0.5, 1)
3   b ~ MvNormal(zeros(degree), 10)
4   # build matrix n*degree
5   t = repeat(mass_std, 1, degree)
6   # exponent its columns
7   t = hcat(map(.^, eachcol(t), 1:degree)...)
8   # calculate product on coefficient's vector
9   μ = a .+ t * b
10
11  log_σ ~ Normal()
12  brain_std ~ MvNormal(μ, exp(log_σ))
13 end

```

	variable	mean	min	median	max	n
1	:a	0.545904	-1.25153	0.532392	1.72525	0
2	Symbol("b[1]")	0.35406	-1.29638	0.363896	1.9635	0
3	Symbol("b[2]")	0.0248361	-2.06992	0.0343043	1.93808	0
4	Symbol("b[3]")	-0.169126	-1.73209	-0.170413	1.42126	0
5	:log_σ	-1.26632	-2.14229	-1.29211	0.937954	0

```

1 begin
2   @time m7_3_ch = sample(model_m7_n(d.mass_std,
3   d.brain_std, degree=3), NUTS(), 1000)
4   @time m7_3 = DataFrame(m7_3_ch)
5   describe(m7_3)
end

100%

```

① Found initial step size
 $\epsilon: 0.2$

Selection 5.452340 seconds (5.31 M allocations: 495.050 MiB, 3.64% gc time, 58.82% compilation time)
0.000095 seconds (122 allocations: 86.250 KiB)

	variable	mean	min	median	max	n
1	:a	0.753135	-1.29212	0.758126	2.73901	0
2	Symbol("b[1]")	0.700627	-2.67157	0.633939	3.60469	0
3	Symbol("b[2]")	-0.684327	-4.93303	-0.586031	4.73013	0
4	Symbol("b[3]")	-0.589297	-3.85836	-0.538768	3.62783	0
5	Symbol("b[4]")	0.479884	-3.22847	0.430074	3.56131	0
6	:log_σ	-1.14958	-1.94352	-1.23606	0.812233	0

```

1 begin
2   @time m7_4_ch = sample(model_m7_n(d.mass_std,
3   d.brain_std, degree=4), NUTS(), 1000)
4   @time m7_4 = DataFrame(m7_4_ch)
5   describe(m7_4)
end

100%

```

① Found initial step size
 $\epsilon: 0.4$

Selection 5.442058 seconds (7.28 M allocations: 865.526 MiB, 4.55% gc time, 33.67% compilation time)
0.000106 seconds (139 allocations: 102.828 KiB)

variable	mean	min	median	max
1 :a	0.807506	-1.39643	0.813605	3.64389
2 Symbol("b[1]")	1.36051	-5.80807	1.42124	7.8659
3 Symbol("b[2]")	-0.448937	-9.47283	-0.336753	7.13566
4 Symbol("b[3]")	-2.11072	-16.9476	-2.28918	14.9686
5 Symbol("b[4]")	0.0860457	-7.38378	-0.0594394	7.71821
6 Symbol("b[5]")	0.779426	-8.71103	0.887674	6.77843
7 :log_σ	-1.02331	-1.89414	-1.07645	0.936759

```

1 begin
2   m7_5_ch = sample(model_m7_n(d.mass_std, d.brain_std,
3   degree=5), NUTS(), 1000)
4   m7_5 = DataFrame(m7_5_ch)
5   describe(m7_5)
end

100%
Selection deleted
ε: 0.025

```

Code 7.9: 6-degree polynomial

```

model_m7_6 (generic function with 2 methods)
1 @model function model_m7_6(mass_std, brain_std)
2   a ~ Normal(0.5, 1)
3   b ~ MvNormal(zeros(6), 10)
4   μ = @. a + b[1]*mass_std + b[2]*mass_std^2 +
5   b[3]*mass_std^3 +
6   b[4]*mass_std^4 + b[5]*mass_std^5 +
7   b[6]*mass_std^6
     brain_std ~ MvNormal(μ, 0.001)
end

```

	variable	mean	min	median	max
1	:a	0.445913	-0.749119	0.505061	0.525469
2	Symbol("b[1]")	0.773959	-1.34346	0.878036	0.915289
3	Symbol("b[2]")	2.00922	1.59882	1.71208	8.04338
4	Symbol("b[3]")	-0.324458	-0.705828	-0.602745	5.33837
5	Symbol("b[4]")	-3.9359	-12.9266	-3.49367	-3.3188
6	Symbol("b[5]")	-0.532412	-4.17428	-0.353911	-0.283259
7	Symbol("b[6]")	1.83423	1.55421	1.63434	5.90625

```

1 begin
2   m7_6_ch = sample(model_m7_6(d.mass_std, d.brain_std),
3   NUTS(), 1000)
4   m7_6 = DataFrame(m7_6_ch)
5   describe(m7_6)
end

100%
Selection deleted
Selection deleted
ε: 9.765625e-5

```

Code 7.10 Fig 7.3

```

mass_seq =
-1.0087216024247314:0.02454648527479833:1.4213804397803034
1 mass_seq = range(extrema(d.mass_std)...; length=100)

▶ [0.361326, 0.3655, 0.369673, 0.373847, 0.378021, 0.382194, 0.386

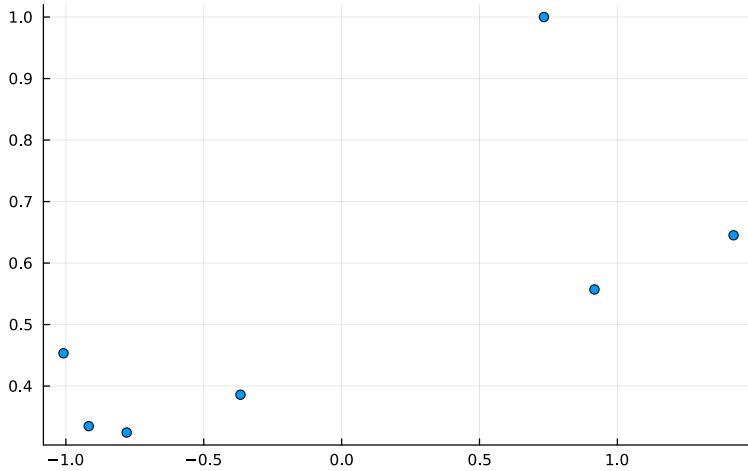
1 begin
2   l = [
3     @. r.a + r.b * mass_seq
4     for r ∈ eachrow(m7_1)
5   ]
6   l = vcat(l'...)
7   μ = mean.(eachcol(l))
8 end

100×2 Matrix{Float64}:
0.115481  0.600084
0.122292  0.59868
0.128817  0.599644
0.135562  0.600578
0.143339  0.602665
0.148106  0.606391
0.153797  0.609951
⋮
0.453709  1.04696
0.453461  1.05329
0.453229  1.0602
0.45253   1.06744
0.452237  1.07638
0.452994  1.08648

1 begin
2   ci = PI.(eachcol(l))
3   ci = vcat(ci'...)
4 end

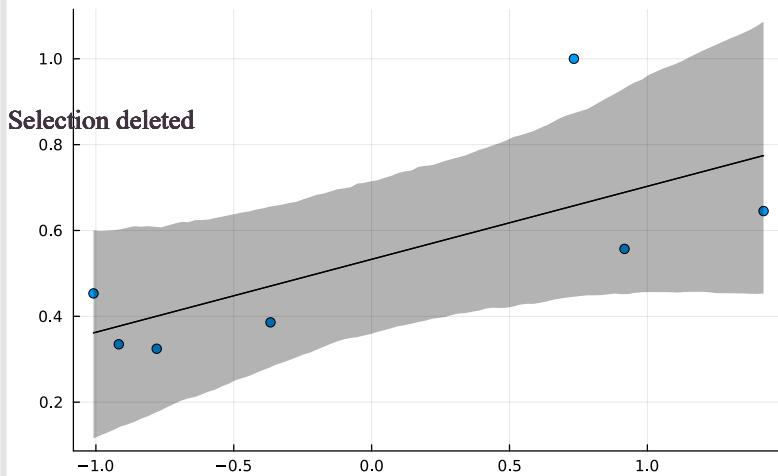
```

1: $R^2 = 0.454$



```
1 scatter(d.mass_std, d.brain_std; title="1: R2 =  
$(round(R2_is_bad(m7_1); digits=3))")
```

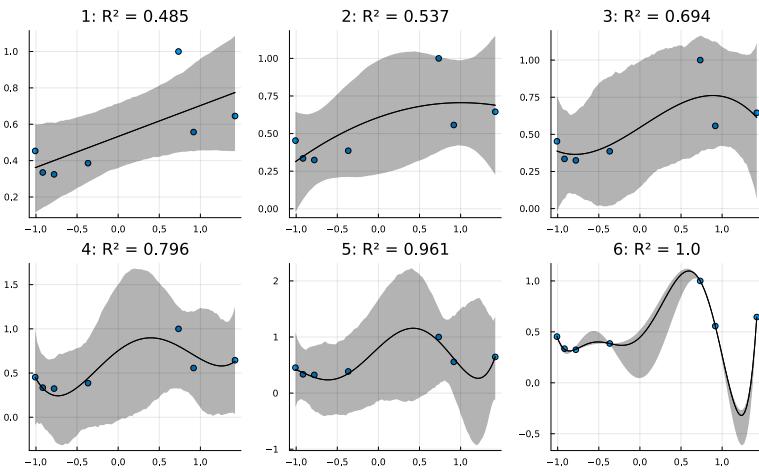
1: $R^2 = 0.454$



```
1 plot!(mass_seq, [μ μ]; fillrange=ci, c=:black,  
fillalpha=0.3)
```

Reimplemented the brain_plot function to check my results.

```
1 function brain_plot(df; sigma=missing)  
2     degree = ncol(df[!,r"b"])  
3     # build mass_seq*degree matrix, with each col  
4     exponentiated to col's index  
5     t = repeat(mass_seq, 1, degree)  
6     t = hcat(map(.^, eachcol(t), 1:degree)...)  
7     l = [  
8         r.a .+ t * collect(r[r"b"])  
9         for r ∈ eachrow(df)  
10    ]  
11    l = vcat(l')...  
12    μ = mean.(eachcol(l))  
13    ci = PI.(eachcol(l))  
14    ci = vcat(ci')...  
15  
16    r2 = round(R2_is_bad(df, sigma=sigma); digits=3)  
17    scatter(d.mass_std, d.brain_std; title="$degree: R2 =  
18 $r2")  
19    plot!(mass_seq, [μ μ]; fillrange=ci, c=:black,  
20          fillalpha=0.3)  
21 end;
```



```

1 plot(
2   brain_plot(m7_1),
3   brain_plot(m7_2),
4   brain_plot(m7_3),
5   brain_plot(m7_4),
6   brain_plot(m7_5),
7   brain_plot(m7_6, sigma=0.001);
8   size=(1000, 600)
9 )

```

Selection deleted Code 7.11

```

1 i = 3;

1 d_minus_i = d[setdiff(1:end,i),:];

brain_loo_plot (generic function with 1 method)
1 function brain_loo_plot(model, data; title::String)
2   (a, b) = extrema(data.brain_std)
3   p = scatter(data.mass_std, data.brain_std; title=title,
4   ylim=(a-0.1, b+0.1))
5   mass_seq = range(extrema(data.mass_std)...; length=100)

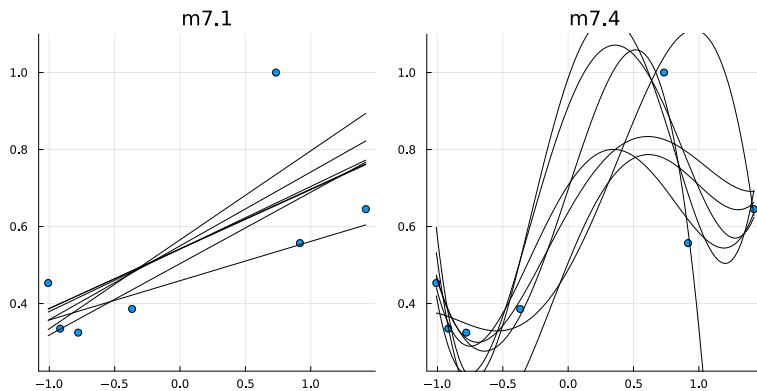
6
7   for i ∈ 1:nrow(data)
8     d_minus_i = data[setdiff(1:end,i),:]
9     df = DataFrame(sample(
10       model(d_minus_i.mass_std,
11       d_minus_i.brain_std),
12       NUTS(),
13       1000))
14
15     degree = ncol(df[!,r"b"])
16
17     # build mass_seq*degree matrix, with each col
18     exponentiated to col's index
19
20     t = repeat(mass_seq, 1, degree)
21     t = hcat(map(.^, eachcol(t), 1:degree)...)
22     l = [
23       r.a .+ t * collect(r[r"b"])
24       for r ∈ eachrow(df)
25     ]
26     l = vcat(l...)
27     μ = mean.(eachcol(l))
28     plot!(mass_seq, μ; c=:black)
end
p
end

```

```

model_m7_4 = #15 (generic function with 1 method)
1 model_m7_4 = (mass, brain) -> model_m7_n(mass, brain,
degree=4)

```



```

1 plot(
2   brain_loo_plot(model_m7_1, d, title="m7.1"),
3   brain_loo_plot(model_m7_4, d, title="m7.4");
4   size=(800, 400)
5 )

```

100%

① Found initial step size
ε: 0.05

100%

② Found initial step size
ε: 0.0125

Selection deleted

100%

① Found initial step size
ε: 0.8

100%

① Found initial step size
ε: 0.025

100%

① Found initial step size
ε: 1.6

100%

① Found initial step size
ε: 0.2

100%

① Found initial step size
ε: 0.2

100%

① Found initial step size
ε: 0.05

100%

7.2 Entropy and accuracy.

Code 7.12 Information and KL Divergence

```
1 p = [0.3, 0.7];
```

```
calc_entropy_given_p (generic function with 1 method)
```

```
1 function calc_entropy_given_p(p)
2     -sum(p .* log.(p))
3 end
```

```
0.6108643020548935
```

```
1 @time calc_entropy_given_p(p)
```

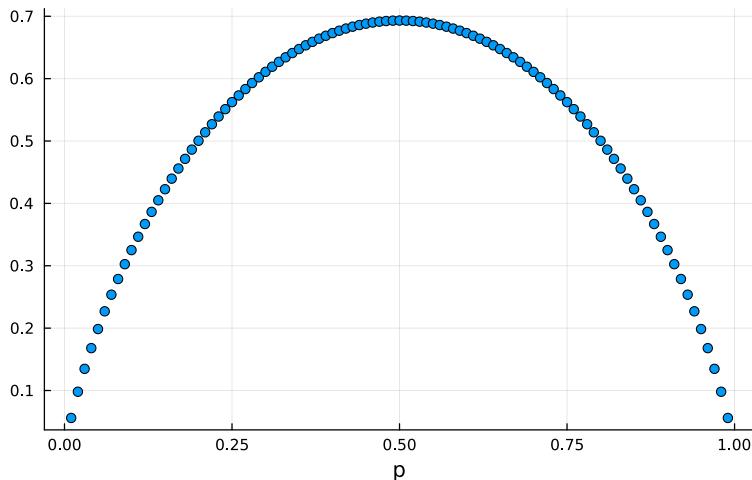
```
0.000002 seconds (1 allocation: 80 bytes)
```

```
plot_entropy_for_binomial_varying_p (generic function with 1 method)
```

```
1 function plot_entropy_for_binomial_varying_p()
2     let p0_vec = 0:0.01:1;
3         entropy_vec = similar(p0_vec);
4         for i ∈ 1:length(p0_vec)
5             entropy_vec[i] =
6                 calc_entropy_given_p([p0_vec[i], 1-p0_vec[i]])
7         end
8         scatter(p0_vec, entropy_vec; title="Entropy for a
9             binomial(p)", xlabel="p");
end
end
```

Selection deleted

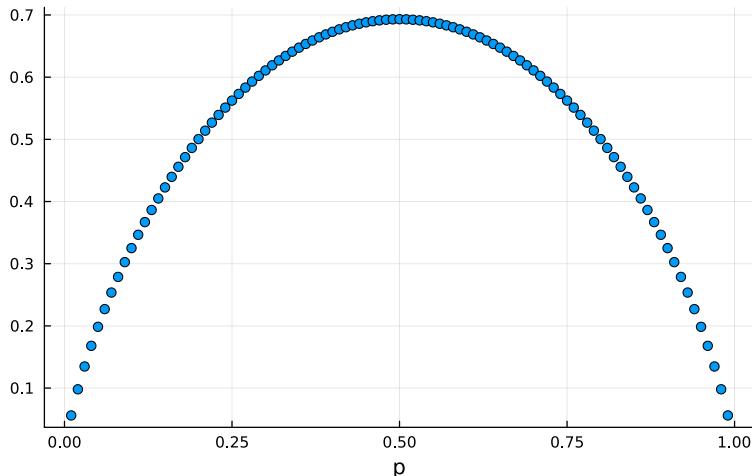
Entropy for a binomial(p)



```
1 @time plot_entropy_for_binomial_varying_p()
```

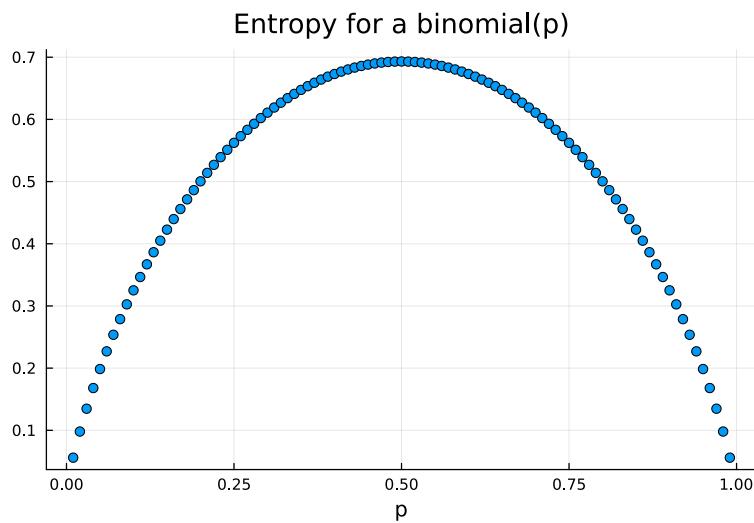
```
0.000806 seconds (701 allocations: 64.781 KiB)
```

Entropy for a binomial(p)



```
1 @time plot_entropy_for_binomial_varying_p()
```

```
0.000706 seconds (701 allocations: 64.781 KiB)
```



```
1 plot_entropy_for_binomial_varying_p()
```

7.13 lppd: Log-Pointwise-Predictive-Density

```
▶ [0.360622, 0.386067, 0.314472, 0.37288, 0.227718, 0.180141, -0.5
Selection deleted
1 lppd(m7_1, (r,x)->Normal(r.a + r.b*x, exp(r.log_σ)),
d.mass_std, d.brain_std)
```

7.14 Run lppd manually

```
▶ [0.360622, 0.386067, 0.314472, 0.37288, 0.227718, 0.180141, -0.5
1 [
2   begin
3     s = [
4       logpdf(Normal(r.a + r.b * x, exp(r.log_σ)), y)
5       for r ∈ eachrow(m7_1)
6     ]
7     # average log likelihood: log(Σi expis/n)
8     # average all n samplings. i is sample i.
9     logsumexp(s) - log(length(s))
10    end
11    for (x, y) ∈ zip(d.mass_std, d.brain_std)
12  ]
```

7.15 lppd function

```
1 # it could be implemented in a generic way, but I'm too lazy
2 df_funcs = [
3     (m7_1, (r, x) -> Normal(r.a + r.b*x, exp(r.log_σ))),
4     (m7_2, (r, x) -> Normal(r.a + r."b[1]" * x +
5     r."b[2]"*x^2, exp(r.log_σ))),
6     (m7_3, (r, x) -> Normal(r.a + r."b[1]" * x +
7     r."b[2]"*x^2 + r."b[3]"*x^3, exp(r.log_σ))),
8     (m7_4, (r, x) -> Normal(r.a + r."b[1]" * x +
9     r."b[2]"*x^2 + r."b[3]"*x^3 +
10    r."b[4]"*x^4,
11    exp(r.log_σ))),
12     (m7_5, (r, x) -> Normal(r.a + r."b[1]" * x +
13     r."b[2]"*x^2 + r."b[3]"*x^3 +
14     r."b[4]"*x^4 +
15     r."b[5]"*x^5, exp(r.log_σ))),
16     (m7_6, (r, x) -> Normal(r.a + r."b[1]" * x +
17     r."b[2]"*x^2 + r."b[3]"*x^3 +
18     r."b[4]"*x^4 +
19     r."b[5]"*x^5 + r."b[6]"*x^6, 0.001)),
20 ];
```

```
► [1.31136, 0.735447, 0.855293, 0.305587, -0.299332, 38.785]
```

```
1 Selection deleted
2 sum(lppd(df, f, d.mass_std, d.brain_std))
3 for (df, f) in df_funcs
4 ]
```

7.16 Define multiple functions

```
m7_sim (generic function with 2 methods)
1 @model function m7_sim(x, y; degree::Int=2)
2     beta ~ MvNormal(zeros(degree), 1)
3     μ = x * beta
4     y ~ MvNormal(μ, 1)
5 end
```

get_lppd

- Calculate sum(-2*lppd) from sampled params (b), an x matrix and target y values.
- call SR.lppd() function.
- Arguments:
 - m_df: data frame of model coefficients.
 - xseq: a matrix of x, each row is one sample.
- yseq: a vector of response variable y.
- Returned values
 - lppd of all samples.

```
1 """
2 - Calculate sum(-2*lppd) from sampled params (b), an x
3 matrix and target y values.
4
5 $(SIGNATURES)
6 - call SR.lppd() function.
Selection deleted
8
9 - Arguments:
10 - m_df: data frame of model coefficients.
11 - xseq: a matrix of x, each row is one sample.
12 - yseq: a vector of response variable y.
13
14 - Returned values
15 - lppd of all samples.
16
17 """
18 function get_lppd(m_df, xseq, yseq)
19     t = DataFrame(:b => collect(eachrow(Matrix(m_df))))
20     -2*sum(StatisticalRethinking.lppd(t, (r, x) =>
21 Normal(r.b'*x, 1), eachrow(xseq), yseq))
end
```

calc_train_test

```
1 """
2 $(SIGNATURES)
3 """
4 function calc_train_test(N, k; count=100)
5     trn_v, tst_v = [], []
6     for _ in 1:count
7         # method sim_train_test from StatisticalRethinking
8         # just simulates the data to be fitted by the model
9         y, x_train, x_test = sim_train_test(N=N, K=k)
10
11         estim = optimize(m7_sim(x_train, y,
12             degree=max(2,k)), MAP())
13         m7_2 = DataFrame(sample(estim, 1000))
14         # commented out is the MCMC way of estimation
15         # instead of MAP
16         # m_chain = sample(m7_sim(x_train, y,
17         # degree=max(2,k)), NUTS(), 1000)
18         # m7_2 = DataFrame(m_chain)
19         t1 = get_lppd(m7_2, x_train, y)
20         t2 = get_lppd(m7_2, x_test, y)
21         push!(trn_v, t1)
22         push!(tst_v, t2)
23     end
24     (mean_and_std(trn_v), mean_and_std(tst_v))
end
```

7.17 Multi-threading

```
TaskFailedException
nested task error: MethodError: no method matching
sample(::TuringOptimExt.ModeResult{NamedArrays.NamedVec
tor{Float64}, Vector{Float64}},
Tuple{OrderedCollections.OrderedDict{Symbol, Int64}}),
Optim.MultivariateOptimizationResults{Optim.LBFGS{Nothi
ng, LineSearches.InitialStatic{Float64},
LineSearches.HagerZhang{Float64, Base.RefValue{Bool}}},
Optim.var"#19#21"}, Vector{Float64}, Float64, Float64,
Vector{Optim.OptimizationState{Float64},
Optim.LBFGS{Nothing,
LineSearches.InitialStatic{Float64},
LineSearches.HagerZhang{Float64, Base.RefValue{Bool}}},
Optim.var"#19#21"}}, Bool,
@NamedTuple{f_limit_reached::Bool,
g_limit_reached::Bool, h_limit_reached::Bool,
time_limit::Bool, callback::Bool, f_increased::Bool}),
DynamicPPL.LogDensityFunction{DynamicPPL.TypedVarInfo{@
NamedTuple{beta::DynamicPPL.Metadata{Dict{AbstractPPL.V
arName{:beta, Setfield.IdentityLens}}, Int64},
Vector{Distributions.IsoNormal},
Vector{AbstractPPL.VarName{:beta,
Setfield.IdentityLens}}, Vector{Float64},
Vector{Set{DynamicPPL.Selector}}}}, Float64},
DynamicPPL.Model{typeof(Main.var"workspace#9".m7_sim),
(:x, :y), (:degree,), (), Tuple{Matrix{Float64},
Vector{Float64}}, Tuple{Int64},
DynamicPPL.DefaultContext},
Turing.Optimisation.OptimizationContext{DynamicPPL.Def
aultContext}}}, ::Int64)
Closest candidates are:
sample(!Matched::DataFrames.DataFrame, ::Any; replace,
ordered)
@ StatisticalRethinking
~/.julia/packages/StatisticalRethinking/RYYWV/src/sampl
e_dataframe.jl:33
sample(!Matched::MCMCChains.Chains, ::Integer; replace,
ordered)
@ MCMCChains
~/.julia/packages/MCMCChains/zFCJy/src/sampling.jl:19
sample(!Matched::Random.AbstractRNG, ::Any,
!Matched::AbstractMCMC.AbstractSampler, !Matched::Any;
kwargs...)
@ AbstractMCMC
~/.julia/packages/AbstractMCMC/YrmkI/src/logdensityprob
lems.jl:43
...
Stacktrace:
[1] #calc_train_test#16
@
~/src/SR2TuringPluto.jl/notebooks/Chapter_07.jl##=##1e50
2519-f1ed-4c83-ba8d-d96e66269d4e:11 [inlined]
[2] macro expansion
@
~/src/SR2TuringPluto.jl/notebooks/Chapter_07.jl##=##823b
783e-f114-427e-b4b4-f135eec09d86:10 [inlined]
[3]
(:Main.var"workspace#9".var"#400#threadsfor_fun#18"
{Main.var"workspace#9".var"#400#threadsfor_fun#17#19"
{UnitRange{Int64}}})(tid::Int64; onethread::Bool)
@ Main.var"workspace#9"
./threadingconstructs.jl#@##=##823b783e-f114-427e-b4b4-
f135eec09d86:215
```

```

[4] #400#threadsfor_fun
@ ./threadingconstructs.jl#@#==#823b783e-f114-427e-
b4b4-f135eec09d86:182 [inlined]
[5] (::Base.Threads.var"#1#2"
{Main.var"workspace#9".var"#400#threadsfor_fun#18"
{Main.var"workspace#9".var"#400#threadsfor_fun#17#19"
{UnitRange{Int64}}}, Int64))()
@ Base.Threads ./threadingconstructs.jl:154
...and 4 more exceptions.

```

Stack trace

Here is what happened, the most recent locations are first:

1. `threading_run(fun=:Main.var"workspace#9".var"#400#threadsfor_fun#18"`
`{Main.var"workspace#9".var"#400#threadsfor_fun#17#19"`
`{UnitRange{Int64}}}, static::Bool)`
`@ | threadingconstructs.jl:172`
2. `macro expansion @ | This cell: line 220`
3. [Show more...](#)

Selection deleted

```

1 begin
2     k_count = 5
3     k_seq = 1:k_count
4     count = 100
5     trn_20, tst_20 = [], []
6     trn_100, tst_100 = [], []
7
8     Threads.@threads for k in k_seq
9         println("Processing $k with N=20...")
10        t1, t2 = calc_train_test(20, k, count=count)
11        push!(trn_20, t1)
12        push!(tst_20, t2)
13        println("Processing $k with N=100...")
14        t1, t2 = calc_train_test(100, k, count=count)
15        push!(trn_100, t1)
16        push!(tst_100, t2)
17    end
18 end

```

```

Processing 1 with N=20...
Processing 2 with N=20...
Processing 5 with N=20...
Processing 4 with N=20...
Processing 3 with N=20...

```

7.18 Plot the training and testing deviance

```
Failed to show value:  
DivideError: integer division error
```

Stack trace

Here is what happened, the most recent locations are first:

```
1. div @ int.jl:295  
2. div @ div.jl:308  
3. div @ div.jl:353  
4. fld @ div.jl:319  
5. mod @ int.jl:287  
6. mod1 @ operators.jl:830  
7. mod @ range.jl:1508  
8. _cycle(v::Vector{Float64}, idx::Int64) @ utils.jl:196  
9. gr_draw_markers(series::Plots.Series,  
    x::UnitRange{Int64}, y::Vector{Float64}, z::Nothing,  
    clims::Tuple{Float64, Float64}, msize::Int64,  
    strokewidth::Int64) @ gr.jl:1850  
10. gr_draw_markers(series::Plots.Series,  
    x::UnitRange{Int64}, y::Vector{Float64}, z::Nothing,  
    clims::Tuple{Float64, Float64}) @ gr.jl:1839  
11. gr_add_series(sp::Plots.Subplot{Plots.GRBackend},  
    series::Plots.Series) @ gr.jl:1744  
12. gr_display(sp::Plots.Subplot{Plots.GRBackend},  
    w::Measures.AbsoluteLength,  
    h::Measures.AbsoluteLength,  
    vp_canvas::Plots.GRViewport{Float64}) @ gr.jl:964  
13. (::Plots.var"#509#510"{Int64, Int64,  
    Plots.GRViewport{Float64}})  
    (sp::Plots.Subplot{Plots.GRBackend}) @ gr.jl:688  
14. foreach(f::Plots.var"#509#510"{Int64, Int64,  
    Plots.GRViewport{Float64}},  
    itr::Vector{Plots.Subplot}) @ (abstractarray.jl:3097)  
15. gr_display(plt::Plots.Plot{Plots.GRBackend},  
    dpi_factor::Int64) @ gr.jl:688  
16. #554 @ gr.jl:2062  
17. withenv(::Plots.var"#554#555"  
    {Plots.Plot{Plots.GRBackend}, Int64}, ::Pair{String,  
    String}, ::Vararg{Pair{String, String}})  
    @ (env.jl:257)  
18. _show(io::IOContext{IOBuffer},  
    ::MIME{Symbol("image/svg+xml")},  
    plt::Plots.Plot{Plots.GRBackend}) @ gr.jl:2057  
19. #invokelatest#2 @ essentials.jl:892  
20. invokelatest @ essentials.jl:889  
21. show(io::IOContext{IOBuffer},  
    m::MIME{Symbol("image/svg+xml")},  
    plt::Plots.Plot{Plots.GRBackend}) @ output.jl:232
```

```
1 begin  
2     scatter(k_seq, first.(trn_20); yerr=last.(trn_20),  
3     label="train", title="N=20")
```

```
4 scatter!(k_seq .+ .1, first.(tst_20); yerr=last.  
(tst_20), label="test")
```

Selection deleted

```
Failed to show value:  
DivideError: integer division error
```

Stack trace

Here is what happened, the most recent locations are first:

```
1. div @ int.jl:295  
2. div @ div.jl:308  
3. div @ div.jl:353  
4. fld @ div.jl:319  
5. mod @ int.jl:287  
6. mod1 @ operators.jl:830  
7. mod @ range.jl:1508  
8. _cycle(v::Vector{Float64}, idx::Int64) @ utils.jl:196  
9. gr_draw_markers(series::Plots.Series,  
    x::UnitRange{Int64}, y::Vector{Float64}, z::Nothing,  
    clims::Tuple{Float64, Float64}, mszie::Int64,  
    strokewidth::Int64) @ gr.jl:1850  
10. gr_draw_markers(series::Plots.Series,  
Selection deleted  
    x::UnitRange{Int64}, y::Vector{Float64}, z::Nothing,  
    clims::Tuple{Float64, Float64}) @ gr.jl:1839  
11. gr_add_series(sp::Plots.Subplot{Plots.GRBackend},  
    series::Plots.Series) @ gr.jl:1744  
12. gr_display(sp::Plots.Subplot{Plots.GRBackend},  
    w::Measures.AbsoluteLength,  
    h::Measures.AbsoluteLength,  
    vp_canvas::Plots.GRViewport{Float64}) @ gr.jl:964  
13. (::Plots.var"#509#510"{Int64, Int64,  
    Plots.GRViewport{Float64}})  
    (sp::Plots.Subplot{Plots.GRBackend}) @ gr.jl:688  
14. foreach(f::Plots.var"#509#510"{Int64, Int64,  
    Plots.GRViewport{Float64}},  
    itr::Vector{Plots.Subplot}) @ abstractarray.jl:3097  
15. gr_display(plt::Plots.Plot{Plots.GRBackend},  
    dpi_factor::Int64) @ gr.jl:688  
16. #554 @ gr.jl:2062  
17. withenv(::Plots.var"#554#555"  
    {Plots.Plot{Plots.GRBackend}, Int64}, ::Pair{String,  
    String}, ::Vararg{Pair{String, String}})  
    @ env.jl:257  
18. _show(io::IOContext{IOBuffer},  
    ::MIME{Symbol("image/svg+xml")},  
    plt::Plots.Plot{Plots.GRBackend}) @ gr.jl:2057  
19. #invokelatest#2 @ essentials.jl:892  
20. invokelatest @ essentials.jl:889  
21. show(io::IOContext{IOBuffer},  
    m::MIME{Symbol("image/svg+xml")},  
    plt::Plots.Plot{Plots.GRBackend}) @ output.jl:232
```

```
1 begin  
2     scatter(k_seq, first.(trn_100); yerr=last.(trn_100),  
3             label="train", title="N=100")  
4     scatter!(k_seq .+ .1, first.(tst_100); yerr=last.  
        (tst_100), label="test")  
end
```

7.3 Golem taming: regularization

No code pieces in this section

7.4 Predicting predictive accuracy

7.19 Monte Carlo of a Bayesian Linear Model

d_cars =	speed	dist
	1 4	2
	2 4	10
	3 7	4
Selection deleted	4 7	22
	5 8	16
	6 9	10
	7 10	18
	8 10	26
	9 10	34
	10 11	17
	: more	
	50 25	85

```
1 d_cars = DataFrame(CSV.File("data/cars.csv", drop= ["Column1"]))
```

```
► (50, 2)
```

```
1 size(d_cars)
```

variable	mean	min	median	max	nmissing	eltype
1 :speed	15.4	4	15.0	25	0	Int64
2 :dist	42.98	2	36.0	120	0	Int64

```
1 describe(d_cars)
```

```
► [5.28764, 25.7694]
```

```
1 std.(eachcol(d_cars))
```

	a	b	σ
1	-13.5798	3.81361	13.1482
2	-22.8556	4.48706	12.5707
3	-22.8556	4.48706	12.5707
4	-22.0897	3.9545	14.1456
5	-19.7891	4.11428	14.4395
6	-23.3199	4.35198	15.4593
7	-16.7861	3.86095	12.0726
8	-18.3343	4.0153	14.9624
9	-19.8632	4.1205	14.6075
10	-20.4602	4.07331	14.5578
: more			
1000	-14.6604	3.71711	13.564

```

1 begin
2 @model function model_m(speed, dist)
3     a ~ Normal(0, 100)
4     b ~ Normal(0, 10)
5     μ = @. a + b * speed
6     σ ~ Exponential(1)
7     dist ~ MvNormal(μ, σ)
8 end
9
10 Random.seed!(17)
11 @time m_cars_ch = sample(model_m(d_cars.speed,
12 d_cars.dist), NUTS(), 1000)
13 m_cars_df = DataFrame(m_cars_ch);
end

```

100%

ⓘ Found initial step size
 $\epsilon: 0.003125$

4.633270 seconds (4.28 M allocations: 394.747 Mi B, 3.92% gc time, 64.14% compilation time)

	variable	mean	min	median	max	nmissing
1	:a	-17.3265	-35.5831	-17.2642	0.299773	0
2	:b	3.91677	2.84743	3.92755	5.00309	0
3	:σ	13.7938	10.0242	13.7182	18.9374	0

1 describe(m_cars_df)

▶ [5.80003, 0.357458, 1.17729]

1 std.(eachcol(m_cars_df))

7.20 logprob

```
1000x50 Matrix{Float64}:
-3.49553 -3.69669 -3.73555 -3.72352 ... -4.06641 -4.15059
-3.60128 -4.15347 -3.51592 -4.02238 -3.6128 -3.66131
-3.60128 -4.15347 -3.51592 -4.02238 -3.6128 -3.66131
-3.73931 -4.22994 -3.57467 -4.24109 -4.48774 -4.5861
-3.65709 -4.01515 -3.64912 -3.99351 -3.99709 -4.06206
-3.78812 -4.18686 -3.67783 -4.11889 ... -3.90446 -3.95205
-3.4482 -3.85121 -3.54348 -3.88427 -4.30168 -4.41574 ...
: ...
-3.56572 -3.87583 -3.65045 -3.88701 -4.09757 -4.17747
-3.69113 -3.74813 -3.9188 -3.78145 ... -4.22037 -4.28792
-3.72779 -3.8742 -3.84009 -3.9114 -4.32572 -4.39368
-3.50471 -3.72501 -3.68585 -3.78668 -4.47811 -4.58639
-3.59089 -3.9929 -3.55534 -4.08819 -5.0411 -5.17345 ...
-3.53509 -3.78693 -3.67355 -3.83406 -4.35388 -4.45144
```

```
1 begin
2     fun = (r, (x, y)) -> normlogpdf(r.a + r.b * x, r.σ, y)
3     @time lp = StatisticalRethinking.link(m_cars_df, fun,
4     zip(d_cars.speed, d_cars.dist))
5     lp = hcat(lp...);
end
```

0.022135 seconds (523.86 k allocations: 8.390 Mi B)

Selection deleted

7.21 lppd

▶ [-3.6192, -3.91929, -3.66225, -3.92798, -3.57279, -3.72152, -3.5

```
1 begin
2     n_samples, n_cases = size(lp)
3     lppd_vals = [
4         logsumexp(c) - log(n_samples)
5         for c in eachcol(lp)
6     ];
7
8     ## if only lppd were needed, we can calculate it with
9     # lppd_vals = lppd(m_df, (r, x) -> Normal(r.a + r.b *
10     x, r.σ), d.speed, d.dist)
end
```

7.22 penalty of WAIC

```
pWAIC = ▶Float64[  
    1: 0.0181654  
    2: 0.081476  
    3: 0.0197703  
    4: 0.0520469  
    5: 0.00931132  
    6: 0.0193994  
    7: 0.00966674  
    8: 0.0100506  
    9: 0.0310315  
   10: 0.0154818  
   11: 0.00811483  
   12: 0.0346919  
   13: 0.0147627  
   14: 0.00953882  
   15: 0.0076483  
   16: 0.0102063  
   17: 0.00746253  
   18: 0.00746253  
   19: 0.018783  
   20: 0.0142461  
    : more  
  41: 0.0167242  
  42: 0.0100626  
  43: 0.00867662  
  44: 0.00947696  
  45: 0.117229  
  46: 0.0219985  
  47: 0.088327  
  48: 0.100576  
  49: 1.26813  
  50: 0.0173583  
]  
1 pWAIC = [  
2   StatisticalRethinking.var2(c)  
3   for c in eachcol(lp)  
4 ]
```

7.23 WAIC

```
420.926630596472  
1 -2*(sum(lppd_vals) - sum(pWAIC))
```

7.24 stddev of WAIC

```
16.19714181056546  
1 begin  
2   waic_vec = -2 * (lppd_vals .- pWAIC)  
3   sqrt(n_cases * StatisticalRethinking.var2(waic_vec))  
4 end
```

7.5 Model comparison

- Data and models from chapter 6

	a	bt	σ
1	1.32535	0.149351	1.75975
2	1.32833	0.133065	1.95495
3	1.31844	0.129862	1.88123
4	1.31068	0.198652	1.63951
5	1.31332	0.199151	1.69585
6	1.35276	0.136793	1.84516
7	1.30178	0.201123	2.01347
8	1.37402	0.098118	1.79777
9	1.36072	0.113754	1.79714
10	1.34972	0.140805	1.67981
⋮ more			
1000	1.34246	0.147348	1.90526

```

1 begin
2 begin
3     Random.seed!(70)
4         # number of plants
5     N = 100
6     h0 = rand(Normal(10, 2), N)
7     treatment = repeat(0:1, inner=div(N, 2))
8     fungus = [rand(Binomial(1, 0.5 - treat*0.4)) for
9     treat in treatment]
10    h1 = h0 .+ rand(MvNormal(5 .- 3 .* fungus, 1))
11
12    d_fungus = DataFrame(:h0 => h0, :h1 => h1,
13    :treatment => treatment, :fungus => fungus)
14
15    @model function model_m6_6(h0, h1)
16        p ~ LogNormal(0, 0.25)
17        σ ~ Exponential(1)
18        μ = h0 .* p
19        h1 ~ MvNormal(μ, σ)
20    end
21
22    @time m6_6 = sample(model_m6_6(d_fungus.h0,
23    d_fungus.h1), NUTS(), 1000)
24    m6_6_df = DataFrame(m6_6)
25
26    @model function model_m6_7(h0, treatment, fungus,
27    h1)
28        a ~ LogNormal(0, 0.2)
29        bt ~ Normal(0, 0.5)
30        bf ~ Normal(0, 0.5)
31        σ ~ Exponential(1)
32        p = @. a + bt*treatment + bf*fungus
33        μ = h0 .* p
34        h1 ~ MvNormal(μ, σ)
35    end
36
37    @time m6_7 = sample(model_m6_7(d_fungus.h0,
38    d_fungus.treatment, d_fungus.fungus, d_fungus.h1),
39    NUTS(), 1000)
40    m6_7_df = DataFrame(m6_7)
41
42    @model function model_m6_8(h0, treatment, h1)
43        a ~ LogNormal(0, 0.2)
44        bt ~ Normal(0, 0.5)
45        σ ~ Exponential(1)
46        p = @. a + bt*treatment
47        μ = h0 .* p
48        h1 ~ MvNormal(μ, σ)

```

```

46      end
47
48      @time m6_8 = sample(model_m6_8(d_fungus.h0,
49          d_fungus.treatment, d_fungus.h1), NUTS(), 1000)
50      m6_8_df = DataFrame(m6_8);
51
52  end

```

100%

- ① Found initial step size
ε: 0.025
- 100%
- ① Found initial step size
ε: 0.0125
- 100%
- ① Found initial step size
ε: 0.003125

3.593926 seconds (3.03 M allocations: 245.770 MiB, 72.44% compilation time)
 5.777817 seconds (4.12 M allocations: 488.726 MiB, 3.72% gc time, 66.40% compilation time)
 4.867473 seconds (3.62 M allocations: 366.757 MiB, 2.19% gc time, 68.24% compilation time)

Selection deleted

7.25 WAIC for m6_7 (include both treatment and fungus)

```

▶ (WAIC = 334.209, lppd = -163.395, penalty = 3.70913, std_err = 1
◀
1 begin
2     @time calc_normlogpdf = (r, (x,bt,bf,y)) ->
3         normlogpdf(x*(r.a + r.bt*bt + r.bf*bf), r.σ, y)
4
5     # log likelihood calculation
6     @time ll = StatisticalRethinking.link(m6_7_df,
7         calc_normlogpdf,
8         zip(d_fungus.h0, d_fungus.treatment,
9             d_fungus.fungus, d_fungus.h1));
10    ll = hcat(ll...);
11    @show size(ll)
12    @time psis.waic(ll)
13
end

```

0.000021 seconds
 0.054654 seconds (1.70 M allocations: 26.684 MiB)
 size(ll) = (1000, 100)
 0.001115 seconds (119 allocations: 1.544 MiB)

7.26 Compare WAIC of m6.6 - m6.8

	models	WAIC	lppd	SE	dWAIC	dSE	pWAIC
1	"m7"	334.2	326.79	13.03	0.0	0.0	3.71
2	"m8"	401.7	394.89	17.42	67.5	14.47	3.42
3	"m6"	418.7	415.09	13.23	84.5	12.27	1.82

```
1 begin
2   calc_normlogpdf_1 = (r, (x,y)) -> normlogpdf(x*r.p,
3   r.o, y)
4   m6_ll = StatisticalRethinking.link(m6_6_df,
5   calc_normlogpdf_1, zip(d_fungus.h0, d_fungus.h1));
6   m6_ll = hcat(m6_ll...);
7
8   calc_normlogpdf_3 = (r, (x,bt,bf,y)) -> normlogpdf(x*
9   (r.a + r.bt*bt + r.bf*bf), r.o, y)
10  m7_ll = StatisticalRethinking.link(m6_7_df,
11  calc_normlogpdf_3, zip(d_fungus.h0, d_fungus.treatment,
12  d_fungus.fungus, d_fungus.h1));
13  m7_ll = hcat(m7_ll...);
14
15  calc_normlogpdf_2 = (r, (x,bt,y)) -> normlogpdf(x*(r.a
16  + r.bt*bt), r.o, y)
17  m8_ll = StatisticalRethinking.link(m6_8_df,
18  calc_normlogpdf_2, zip(d_fungus.h0, d_fungus.treatment,
19  d_fungus.h1));
20  m8_ll = hcat(m8_ll...);

21 compare([m6_ll, m7_ll, m8_ll], :waic, mnames=["m6",
22 "m7", "m8"])
end
```

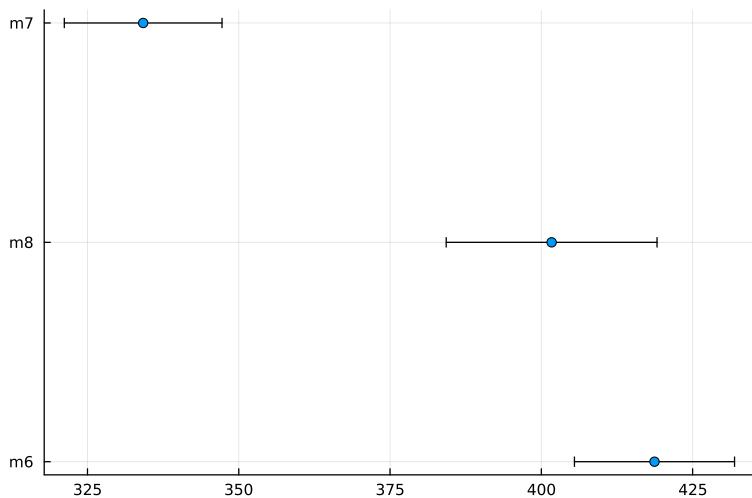
7.27 stddev of WAIC delta between m6.7 and m6.8

```
14.397840569160383
1 begin
2   waic_m6_7 = waic(m7_ll, pointwise=true).WAIC
3   waic_m6_8 = waic(m8_ll, pointwise=true).WAIC
4   n = length(waic_m6_7)
5   diff_m6_78 = waic_m6_7 - waic_m6_8
6   sqrt(n*StatisticalRethinking.var2(diff_m6_78))
7 end
```

7.28 99% confidence interval of dWAIC

```
► [12.96, 67.04]
1 40.0 .+ [-1, 1]*10.4*2.6
```

7.29 plot WAIC/deviance of 3 different models



```
1 begin
2   dw = compare([m6_ll, m7_ll, m8_ll], :waic, mnames=
3     ["m6", "m7", "m8"])
Selection deleted
4   scatter(reverse(dw.WAIC), reverse(dw.models);
5   xerror=reverse(dw.SE))
6 end
```

7.30 stddev of WAIC differences between m6.6 and m6.8

9.896685041881293

```
1 begin
2   waic_m6_6 = waic(m6_ll, pointwise=true).WAIC
3   waic_m6_8_2 = waic(m8_ll, pointwise=true).WAIC
4   diff_m6_68 = waic_m6_6 - waic_m6_8_2
5   sqrt(n*StatisticalRethinking.var2(diff_m6_68))
6 end
```

7.31 stddev of WAIC differences among all models

Current version of StatisticalRethinking.compare doesn't calculate pairwise error. You should use above logic to get values not returned in compare result.

7.32 Fit 3 models of the divorce dataset (Divorce rate ~ Marriage rate, Age at Marriage) and get the corresponding pointwise log-likelihood.

- The divorce datasets contain 50 data points/states.

	a	bA	bM	σ
1	-0.0558592	-0.564404	-0.227424	0.881001
2	0.0269661	-0.586021	0.115961	0.932721
3	0.0941538	-0.600504	-0.0344909	0.864356
4	-0.00379184	-0.637799	-0.128804	0.960596
5	-0.00379184	-0.637799	-0.128804	0.960596
6	0.0439258	-0.404662	-0.104283	0.672458
7	0.0439258	-0.404662	-0.104283	0.672458
8	-0.0231142	-0.478822	-0.0779828	0.795121
9	-0.113931	-0.675914	-0.106931	0.732142
10	-0.115146	-0.408913	0.0772604	0.678882
: more				
1000	0.0493391	-0.774777	-0.0907425	0.834947

```

1 begin
2 Random.seed!(1)
Selection deleted
3 d_divorce =
4 DataFrame(CSV.File("data/WaffleDivorce.csv"))
5 d_divorce[:,D] = standardize(ZScoreTransform,
6 d_divorce.Divorce)
7 d_divorce[:,M] = standardize(ZScoreTransform,
8 d_divorce.Marriage)
9 d_divorce[:,A] = standardize(ZScoreTransform,
10 d_divorce.MedianAgeMarriage)
11 @show size(d_divorce)
12 @model function model_m5_1(A, D)
13     σ ~ Exponential(1)
14     a ~ Normal(0, 0.2)
15     bA ~ Normal(0, 0.5)
16     μ = @. a + bA * A
17     D ~ MvNormal(μ, σ)
18 end
19
20 @time m5_1 = sample(model_m5_1(d_divorce.A,
21 d_divorce.D), NUTS(), 1000)
22 m5_1_df = DataFrame(m5_1)
23
24 @model function model_m5_2(M, D)
25     σ ~ Exponential(1)
26     a ~ Normal(0, 0.2)
27     bM ~ Normal(0, 0.5)
28     μ = @. a + bM * M
29     D ~ MvNormal(μ, σ)
30 end
31
32 @time m5_2 = sample(model_m5_2(d_divorce.M,
33 d_divorce.D), NUTS(), 1000)
34 m5_2_df = DataFrame(m5_2);
35
36 @model function model_m5_3(A, M, D)
37     σ ~ Exponential(1)
38     a ~ Normal(0, 0.2)
39     bA ~ Normal(0, 0.5)
40     bM ~ Normal(0, 0.5)
41     μ = @. a + bA * A + bM * M
42     D ~ MvNormal(μ, σ)
43 end
44
45 @time m5_3 = sample(model_m5_3(d_divorce.A,
46 d_divorce.M, d_divorce.D), NUTS(), 1000)
47 m5_3_df = DataFrame(m5_3);
48
end

```

```
100%  
① Found initial step size  
ε: 0.05
```

```
100%  
① Found initial step size  
ε: 0.025
```

```
100%  
① Found initial step size  
ε: 0.025
```

```
size(d_divorce) = (50, 16)  
 4.255711 seconds (3.26 M allocations: 263.999 MiB,  
 2.36% gc time, 68.55% compilation time)  
 4.245426 seconds (3.25 M allocations: 261.452 MiB,  
 1.43% gc time, 69.63% compilation time)  
 5.417593 seconds (3.74 M allocations: 333.725 MiB,  
 3.13% gc time, 64.58% compilation time)
```

```
1000×50 Matrix{Float64}:  
-2.00667 -2.37024 -0.996025 -2.71984 ... -0.879512 -1.17348  
-1.77528 -1.35202 -0.970049 -1.48071 ... -0.852814 -1.04555  
-1.73167 -1.57019 -0.877906 -1.75229 ... -0.777612 -1.16394  
-1.75853 -1.80834 -1.00913 -1.95755 ... -0.900331 -1.16842  
-1.75853 -1.80834 -1.00913 -1.95755 ... -0.900331 -1.16842  
-2.58925 -2.64059 -0.786789 -3.53794 ... -0.646975 -1.24077  
-2.58925 -2.64059 -0.786789 -3.53794 ... -0.646975 -1.24077  
⋮  
-1.89483 -1.73666 -0.837993 -1.91391 ... -0.690428 -1.09316  
-2.32464 -2.15832 -0.957138 -2.51128 ... -0.74571 -0.961021  
-2.30974 -2.21973 -0.948318 -2.43778 ... -0.701305 -0.934621  
-1.81103 -1.94061 -0.689592 -2.05472 ... -0.562244 -1.30165  
-2.05383 -2.93154 -0.735413 -2.9586 ... -0.531055 -1.34278  
-1.66605 -1.61211 -0.857697 -1.60905 ... -0.738616 -1.11867
```

```
Selection deleted  
1 begin  
2   @time m5_1_ll = StatisticalRethinking.link(m5_1_df,  
3     (r, (x,y)) -> StatsFuncs.normlogpdf(r.a + r.bA * x,  
4       r.σ, y),  
5     zip(d_divorce.A, d_divorce.D));  
6   m5_1_ll = hcat(m5_1_ll...)  
7   @show size(m5_1_ll)  
8  
9   @time m5_2_ll = StatisticalRethinking.link(m5_2_df,  
10    (r, (x,y)) -> StatsFuncs.normlogpdf(r.a + r.bM * x,  
11      r.σ, y),  
12    zip(d_divorce.M, d_divorce.D));  
13   m5_2_ll = hcat(m5_2_ll...)  
14   @show size(m5_2_ll)  
15  
16   @time m5_3_ll = StatisticalRethinking.link(m5_3_df,  
17     (r, (a,m,y)) -> StatsFuncs.normlogpdf(r.a + r.bA * a  
18     + r.bM * m, r.σ, y),  
19     zip(d_divorce.A, d_divorce.M, d_divorce.D));  
20   m5_3_ll = hcat(m5_3_ll...);  
end
```

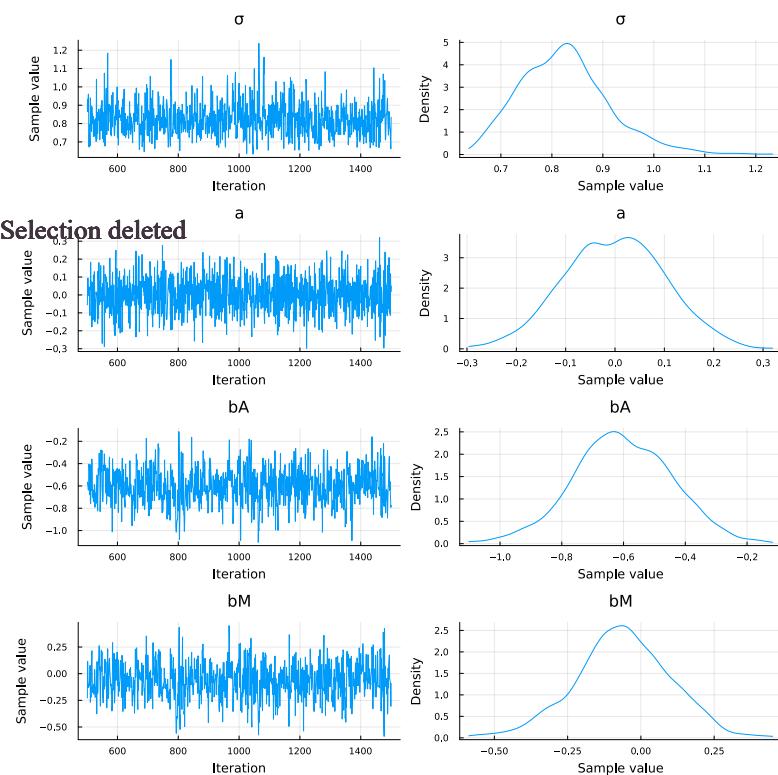
```
0.026356 seconds (623.86 k allocations: 9.916 MiB)  
B)  
size(m5_1_ll) = (1000, 50)  
0.021203 seconds (623.86 k allocations: 9.916 MiB)  
size(m5_2_ll) = (1000, 50)  
0.025481 seconds (848.31 k allocations: 13.342 MiB)
```

	variable	mean	min	median	max	nr
1	:a	-0.000837454	-0.2971	0.00141241	0.320335	0
2	:bA	-0.603134	-1.10509	-0.607978	-0.113036	0
3	:bM	-0.0623091	-0.587797	-0.0600914	0.450345	0
4	:σ	0.825118	0.633363	0.820318	1.23687	0

```
1 describe(m5_3_df)
```

► [0.101411, 0.159148, 0.162949, 0.0882977]

```
1 std.(eachcol(m5_3_df))
```



```
1 plot(m5_3)
```

7.33 Compare 3 models via PSIS: m5_1 is best.

	models	PSIS	lppd	SE	dPSIS	dSE	pPSIS
1	"m5.1"	124.6	118.5	12.26	0.0	0.0	3.28
2	"m5.3"	126.7	118.19	12.28	2.1	0.77	4.6
3	"m5.2"	139.1	133.48	9.61	14.5	8.82	3.03

```
1 compare([m5_1_ll, m5_2_ll, m5_3_ll], :psis, mnames=["m5.1", "m5.2", "m5.3"])
```

7.34 m5_3: Compare pointwise PSIS Pareto k and WAIC penalty

```
► [0.101411, 0.159148, 0.162949, 0.0882977]
```

```
1 begin
2   @show describe(m5_3_df)
3   std.(eachcol(m5_3_df))
4 end
```

Row	variable	mean	min	median
max	nmissing	eltype		
	Symbol	Float64	Float64	Float64
	Float64	Int64	DataType	
1	a	-0.000837454	-0.2971	0.00141241
0.320335		0	Float64	
2	bA	-0.603134	-1.10509	-0.607978
-0.113036		0	Float64	
3	bM	-0.0623091	-0.587797	-0.0600914
0.450345		0	Float64	
4	σ	0.825118	0.633363	0.820318
1.23687		0	Float64	

Selection deleted
`ll_to_psis`

- reshape data to format of ParetoSmooth.psis_loo function
- Arguments
 - `ll`: a [number-of-samplings X number-of-data-points] matrix of log-likelihoods.
- Returned values:
 - a 3D array: [number-of-data-points X number-of-samples X 1].

```
1 """
2 - reshape data to format of ParetoSmooth.psis_loo function
3
4 $(SIGNATURES)
5
6 - Arguments
7   - ll: a [number-of-samplings X number-of-data-points]
8   matrix of log-likelihoods.
9
10 - Returned values:
11   - a 3D array: [number-of-data-points X number-of-samples
12   X 1].
13 """
14 function ll_to_psis(ll::Matrix{<:Real})
15   @show size(ll)
16   t = ll'
17   @show size(t)
18   #Make the dataset 3 dimensional
   collect(reshape(t, size(t)..., 1))
end
```

```
► (WAIC = [4.10301, 4.01367, 1.83596, 4.51732, 1.99058, 2.86319, 2
```

```
1 begin
2     m5_3_t = ll_to_psis(m5_3_ll)
3     @show size(m5_3_t)
4     PSIS_m5_3 = ParetoSmooth.psis_loo(m5_3_t)
5     WAIC_m5_3 = psis.waic(m5_3_ll, pointwise=true)
6 end
```

- ⓘ No source provided for samples; variables are assumed to be from a Markov Chain. If the samples are independent, specify this with keyword argument `source=:other`.

```
size(ll) = (1000, 50)
size(t) = (50, 1000)
size(m5_3_t) = (50, 1000, 1)
```

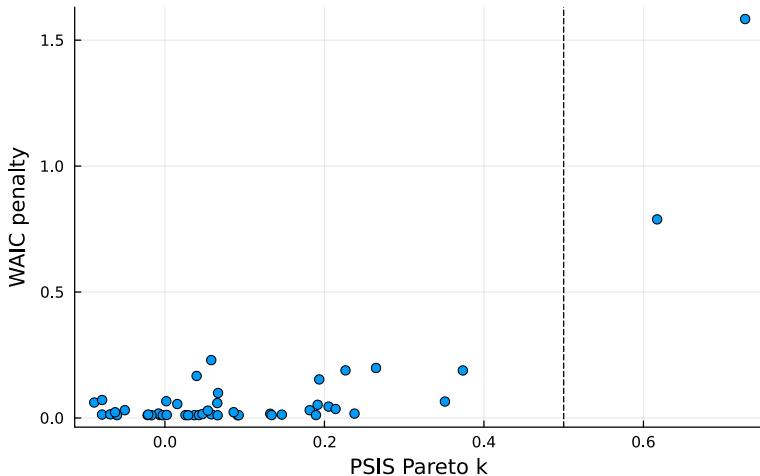
```
1 PSIS_m5_3
```

- ⚠ Some Pareto k values are high (>.7), indicating PSIS has failed to approximate the true distribution.

```
Results of PSIS-LOO-CV with 1000 Monte Carlo samples and 50 data points. Total Monte Carlo SE of 0.14.
```

Selection deleted	total	se_total	mean	se_mean
cv_elpd	-63.83	6.37	-1.28	0.13
naive_lpd	-59.09	4.67	-1.18	0.09
p_eff	4.74	1.82	0.09	0.04

Gaussian model (m5.3)



```
1 begin
2     scatter(PSIS_m5_3.pointwise(:pareto_k),
3             WAIC_m5_3.penalty,
4             xlab="PSIS Pareto k", ylab="WAIC penalty",
5             title="Gaussian model (m5.3)")
6     vline!([0.5], c=:black, s=:dash)
end
```

find_outliers

Find outliers with either Pareto K too large or WAIC penalty too large.

Returned arguments

- a new dataframe given selected rows

```
1 """
2 Find outliers with either Pareto K too large or WAIC
3 penalty too large.
4 $(SIGNATURES)
5
6 Returned arguments
7 - a new dataframe given selected rows
8 """
9 find_outliers(df::DataFrame, pareto_k_vec,
10   waic_penalty_vec::Vector{Float64}; min_pareto_k::Float64=0,
11   min_waic_penalty::Float64=0) = begin
12   state_ind_vec = [i for i in 1:length(pareto_k_vec) if
13     pareto_k_vec[i] >= min_pareto_k ||
14     waic_penalty_vec[i]>=min_waic_penalty]
Selection deleted
15   new_df = DataFrame(d_divorce[state_ind_vec,:],
16     "ParetoK"=>pareto_k_vec[state_ind_vec],
17     "WAIC_Penalty"=>waic_penalty_vec[state_ind_vec])
18
19   # Combine with hcat
20   new_df = DataFrames.hcat(df[state_ind_vec,:],
21     pareto_k_vec[state_ind_vec],
22     waic_penalty_vec[state_ind_vec]; makeunique=true)
23   # Set column names (optional)
24   DataFrames.rename!(new_df, append!(names(df),
25     ["ParetoK", "WAIC_Penalty"]))
26   new_df
end
```

	Location	Loc	Population	MedianAgeMarriage	Marriage
1	"Idaho"	"ID"	1.57	23.2	25.8
2	"Maine"	"ME"	1.33	26.4	13.5

```
1 find_outliers(d_divorce, PSIS_m5_3.pointwise(:pareto_k),
2   WAIC_m5_3.penalty; min_pareto_k=0.4, min_waic_penalty=0.5)
```

7.34.1 m5_3_log_σ_normal: Compare pointwise PSIS Pareto k and WAIC penalty

```
model_m5_3_log_σ (generic function with 2 methods)
1 @model function model_m5_3_log_σ(A, M, D)
2   log_σ ~ Normal(0, 0.2)
3   a ~ Normal(0, 0.2)
4   bA ~ Normal(0, 0.5)
5   bM ~ Normal(0, 0.5)
6   μ = @. a + bA * A + bM * M
7   D ~ MvNormal(μ, exp(log_σ))
8 end
```

	a	bA	bM	log_σ
1	-0.0597823	-0.687137	-0.208746	-0.0589428
2	-0.121128	-0.660067	-0.248804	-0.0354727
3	-0.139087	-0.724421	0.053917	-0.129499
4	0.122024	-0.794938	-0.0262842	-0.194692
5	-0.111765	-0.522158	-0.149208	-0.105211
6	-0.0946616	-0.660986	-0.066957	-0.183706
7	-0.0880742	-0.604822	-0.0846655	-0.231954
8	0.131327	-1.04851	-0.346722	-0.0505851
9	-0.119918	-0.425962	0.361554	-0.294779
10	0.0601405	-0.929196	-0.264513	-0.208219
⋮ more				
1000	-0.0411175	-0.540578	-0.206741	-0.321188

```

1 begin
2 @time m5_3_log_σ = sample(model_m5_3_log_σ(d_divorce.A,
Selection deleted
d_divorce.M, d_divorce.D), NUTS(), 1_000)
3 m5_3_log_σ_df = DataFrame(m5_3_log_σ);
4 end

```

100%

① Found initial step size
ε: 0.2

5.387357 seconds (3.78 M allocations: 340.967 Mi B, 2.72% gc time, 63.30% compilation time)

```

est_summary (generic function with 1 method)
1 est_summary(mcmc_df; σ_ind=4) = begin
2   mcmc_df_sum = describe(mcmc_df)
3   println("exp(median(log_σ estimate) ", exp.
4   (mcmc_df_sum[σ.ind, :median]))
5   println("σ median est is ", median(exp.(mcmc_df[!, :
6   log_σ])))
7   mcmc_df_sum[!, :σ] = std.(eachcol(mcmc_df))
8   @show mcmc_df_sum
9   "stddev of σ is $(std(exp.(mcmc_df[!, :log_σ])))"
end

```

"stddev of σ is 0.08817422802326008"

1 est_summary(m5_3_log_σ_df)

```

exp(median(log_σ estimate) 0.857319310631472
σ median est is 0.8573193134942809
mcmc_df_sum = 4×8 DataFrame
  Row | variable      mean       min       median
  max   | nmissing eltype    σ
  |      Symbol   Float64  Float64  Float64
  |      Int64    DataType  Float64
  └─────────────────────────────────────────────────
    1 | a        -0.00208058 -0.335236  0.000526995
    0.375616 | bA        0.0          0.102158
    2 | bA        -0.599445 -1.09072  -0.604161
    0.0760033 | bM        0.0          0.153298
    3 | bM        -0.0553365 -0.540503 -0.0593282
    0.733242 | log_σ     -0.149939 -0.419509 -0.153945
    0.248442 | log_σ     0.0          0.100376

```

```
► (WAIC = [3.94292, 3.81473, 1.89714, 4.31483, 2.03468, 2.8069, 2.
```

```
1 begin
2   @time m5_3_log_σ_ll = link(m5_3_log_σ_df,
3     (r, (a,m,d)) -> StatsFuns.normlogpdf(r.a + r.bA * a
4     + r.bM * m, exp(r.log_σ), d),
5     zip(d_divorce.A, d_divorce.M, d_divorce.D));
6   m5_3_log_σ_ll = hcat(m5_3_log_σ_ll...);
7   @time PSIS_m5_3_log_σ =
8   ParetoSmooth.psis_loo(ll_to_psis(m5_3_log_σ_ll))
9   @time WAIC_m5_3_log_σ = psis.waic(m5_3_log_σ_ll,
10   pointwise=true)
end
```

- ⓘ No source provided for samples; variables are assumed to be from a Markov Chain. If the samples are independent, specify this with keyword argument `source=:other`.

```
0.184650 seconds (1.08 M allocations: 26.523 MiB, 85.84% compilation time)
size(ll) = (1000, 50)
size(t) = (50, 1000)
0.006581 seconds (1.52 k allocations: 4.180 MiB)
0.000539 seconds (63 allocations: 790.562 KiB)
```

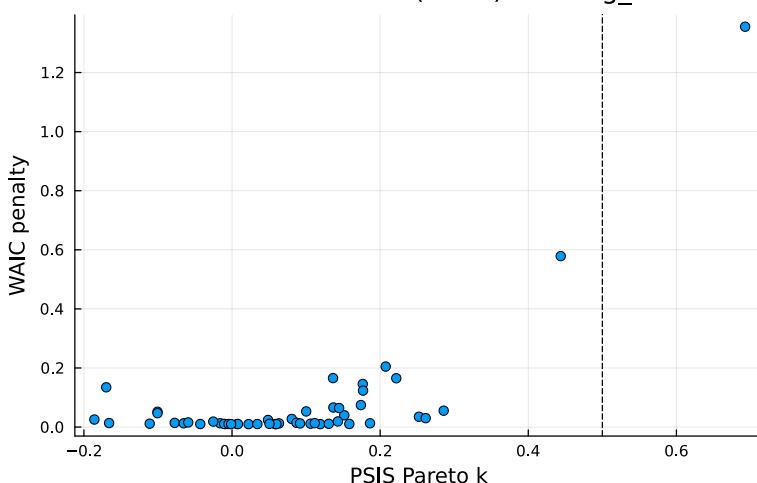
1 PSIS_m5_3_log_σ

- ⓘ Some Pareto k values are slightly high (>.5); some pointwise estimates may be slow to converge or have high variance.

```
Results of PSIS-LOO-CV with 1000 Monte Carlo samples and 50 data points. Total Monte Carlo SE of 0.1.
```

	total	se_total	mean	se_mean
cv_elpd	-63.36	5.67	-1.27	0.11
naive_lpd	-59.50	4.30	-1.19	0.09
p_eff	3.86	1.47	0.08	0.03

Gaussian model (m5.3) with log_σ



```
1 begin
2   scatter(PSIS_m5_3_log_σ.pointwise(:pareto_k),
3     WAIC_m5_3_log_σ.penalty,
4     xlabel="PSIS Pareto k", ylabel="WAIC penalty",
5     title="Gaussian model (m5.3) with log_σ")
6   vline!([0.5], c=:black, s=:dash)
end
```

Location	Loc	Population	MedianAgeMarriage	Marriage
1	"Idaho"	"ID"	1.57	23.2
2	"Maine"	"ME"	1.33	26.4

```
1 find_outliers(d_divorce,
PSIS_m5_3_log_σ.pointwise(:pareto_k),
WAIC_m5_3_log_σ.penalty; min_pareto_k=0.4,
min_waic_penalty=0.3)
```

Base

```
1 @which(names)
```

7.35 m5_3t: MV T-dist df=2, Compare pointwise PSIS Pareto k and WAIC penalty.

Selection deleted

- logNormal for σ is way better than Exponential, manifested by pareto_k.

	a	bA	bM	log_σ
1	0.0431919	-0.413591	-0.0924977	0.164298
2	-0.038226	-0.461917	-0.0751188	-0.287217
3	-0.189934	-0.491807	0.271032	-0.100144
4	-0.0223252	-0.476613	0.1827	-0.166115
5	-0.101952	-0.702216	-0.284762	-0.141744
6	-0.0716474	-0.401141	0.115855	-0.116104
7	0.0323814	-0.604936	-0.125901	-0.151199
8	-0.0525956	-0.902033	-0.311722	0.146498
9	-0.0109573	-0.842162	-0.207502	0.236001
10	0.111475	-0.423923	0.024449	-0.14612
: more				
1000	-0.0964225	-0.588292	-0.035434	0.0609817

```

1 begin
2 @model function model_m5_3t(A, M, D)
3     #σ ~ Exponential(1)
4     # logNormal is way better than Exponential,
5     # manifested by pareto_k.
6     log_σ ~ Normal(0, 0.3)
7     a ~ Normal(0, 0.2)
8     bA ~ Normal(0, 0.5)
9     bM ~ Normal(0, 0.5)
10    μ = @. a + bA * A + bM * M
11    #z = (D.-μ)/σ
12    # The vector . below is optional for NUTS sample.
13    # But psis_loo() requires it.
14    #z ~ TDist(2)
15    #D ~ Distributions.IsoTDist(2, μ, σ)
16    D ~ Distributions.mvtdist(2, μ, exp(log_σ))
17 end
18
19 @time m5_3t = sample(model_m5_3t(d_divorce.A,
d_divorce.M, d_divorce.D), NUTS(), 1000)
m5_3t_df = DataFrame(m5_3t);
end

```

100%

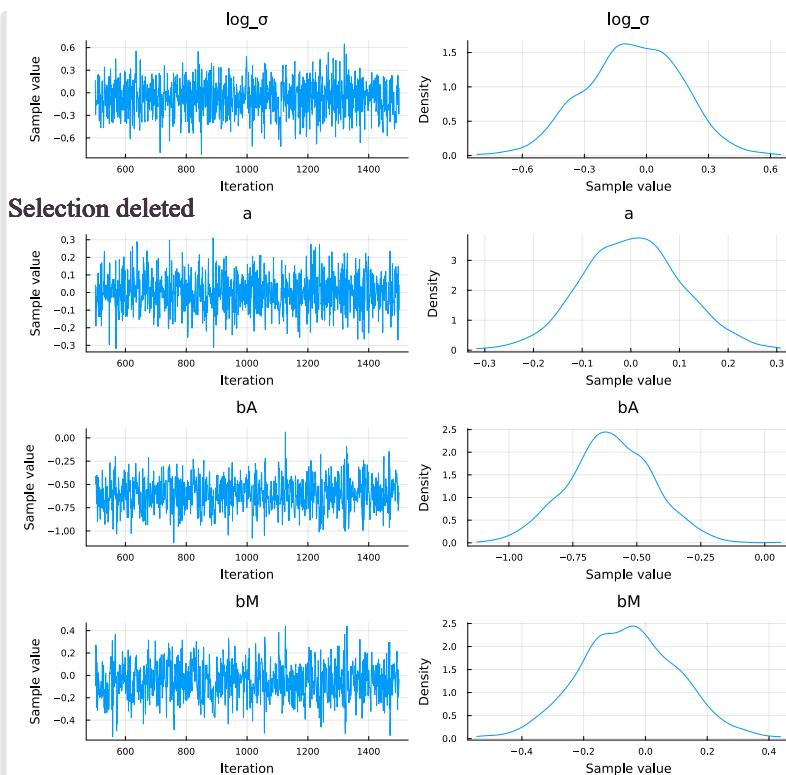
ⓘ Found initial step size
ε: 0.4

⌚ 5.208111 seconds (3.79 M allocations: 335.960 MiB, 2.22% gc time, 67.53% compilation time) ⓘ

```
"stddev of σ is 0.22022987430308238"
```

```
1 est_summary(m5_3t_df)
```

exp(median(log_σ estimate) 0.941349073343889 σ median est is 0.9413490951127395 mcmc_df_sum = 4×8 DataFrame				
Row	variable	mean	min	median
ax	nmissing	Float64	σ	Float64
loat64	Symbol	Int64	DataType	Float64
0.310176	0	0	Float64	0.102272
0.0628147	-0.603098	-1.12604	0	-0.600879
0.440388	0	0	Float64	0.166061
0.649521	-0.0551153	-0.548575	0	-0.0561427
1	a	0.00266562	-0.318649	0.0021977
2	bA	-0.603098	-1.12604	-0.600879
3	bM	-0.0551153	-0.548575	-0.0561427
4	log_σ	-0.0670473	-0.820551	-0.0604412



```
1 plot(m5_3t)
```

```
1 ParetoSmooth.psis_loo(model_m5_3t(d_divorce.A, d_divorce.M,  
d_divorce.D), m5_3t)
```

ⓘ No source provided for samples; variables are assumed to be from a Markov Chain. If the samples are independent, specify this with keyword argument 'source=:other'.

⚠ Some Pareto k values are extremely high (>1). PSIS will not produce consistent estimates.

Results of PSIS-LOO-CV with 1000 Monte Carlo samples and 1 data points. Total Monte Carlo SE of 0.45.				
	total	se_total	mean	se_mean
cv_elpd	-64.19	NaN	-64.19	NaN
naive_lpd	-62.04	NaN	-62.04	NaN
p_eff	2.15	NaN	2.15	NaN

```
► (WAIC = [4.06068, 3.93203, 2.40157, 4.20458, 2.56256, 3.2509, 3.
```

```
1 begin
2
3     m5_3t_ll = StatisticalRethinking.link(m5_3t_df,
4         (r, (a,m,y)) ->
5             Distributions.logpdf(Distributions.mvtdist(2, [r.a
6                 + r.bA * a + r.bM * m], exp(r.log_σ)), [y]),
7                 zip(d_divorce.A, d_divorce.M, d_divorce.D));
8     m5_3t_ll = hcat(m5_3t_ll...);
9
10    m5_3t_t = ll_to_psis(m5_3t_ll)
11    @show size(m5_3t_t)
12    PSIS_m5_3t = ParetoSmooth.psis_loo(m5_3t_t)
13    WAIC_m5_3t = psis.waic(m5_3t_ll, pointwise=true)
14 end
```

- ⓘ No source provided for samples; variables are assumed to be from a Markov Chain. If the samples are independent, specify this with keyword argument `source=:other`.

```
size(ll) = (1000, 50)
size(t) = (50, 1000)
size(m5_3t_t) = (50, 1000, 1)
```

Selection deleted

```
Results of PSIS-LOO-CV with 1000 Monte Carlo samples and 50 data points. Total Monte Carlo SE of 0.057.
```

	total	se_total	mean	se_mean
cv_elpd	-70.67	3.24	-1.41	0.06
naive_elpd	-68.12	3.11	-1.36	0.06
p_eff	2.55	0.20	0.05	0.00

1 PSIS_m5_3t.psis_object

Results of PSIS with 1000 posterior samples and 50 cases.		
pareto_k	ess	sup_ess
0.12	746.61	331.96
0.05	799.58	307.70
0.13	935.96	568.16
0.17	987.88	377.35
0.19	876.83	439.71
0.05	786.54	413.28
0.17	989.91	342.25
0.04	778.72	383.53
0.35	630.51	135.21
0.08	650.16	329.76
0.00	823.99	476.27
0.27	627.98	202.37
0.35	899.14	149.75
0.12	874.61	479.67
0.02	901.77	560.95
0.06	938.39	487.33
0.10	935.95	489.79
0.16	927.72	426.33
-0.03	917.22	623.09
0.08	630.02	164.95
0.05	930.02	478.47
0.02	889.95	449.06
0.09	905.45	480.76
0.20	663.79	199.66
-0.04	859.18	568.05
0.10	791.92	425.28
0.09	645.79	367.49
0.26	657.33	303.14
-0.02	989.06	631.64
0.14	1120.81	348.26
0.06	950.78	503.74
0.26	874.41	260.23
0.05	938.52	486.03
0.18	757.13	252.13
0.05	955.93	501.29
0.11	1174.78	517.91
0.06	949.16	535.75
0.17	755.89	714.84

\emptyset				
(This table has no columns)				
1	1	:cv_elpd	-2.03076	
2	2	:cv_elpd	-1.96683	
3	3	:cv_elpd	-1.20084	
4	4	:cv_elpd	-2.10265	
5	5	:cv_elpd	-1.28145	
6	6	:cv_elpd	-1.6257	
7	7	:cv_elpd	-1.64614	
8	8	:cv_elpd	-1.05047	
9	9	:cv_elpd	-1.17854	
10	10	:cv_elpd	-1.24652	
⋮ more				

```

1 begin
2   @show propertynames(PSIS_m5_3t)
3   @show size(PSIS_m5_3t.pointwise)
4   @show typeof(PSIS_m5_3t.pointwise)
Selection deleted propertynames(PSIS_m5_3t.pointwise)
6   PSIS_m5_3t.pointwise
7 end

```

```

propertynames(PSIS_m5_3t) = (:estimates, :pointwise ⓘ
e, :psis_object, :gmpd, :mcse)
size(PSIS_m5_3t.pointwise) = (50, 5)
typeof(PSIS_m5_3t.pointwise) = AxisKeys.KeyedArray{Flo
at64, 2, NamedDims.NamedDimsArray{(:data, :statistic),
Float64, 2, Matrix{Float64}}}, Tuple{UnitRange{Int64},
Vector{Symbol}}}
propertynames(PSIS_m5_3t.pointwise) = (:data, :statist
ic)

```

▶[:cv_elpd, :naive_lpd, :p_eff, :mcse, :pareto_k]

```

1 begin
2   @show PSIS_m5_3t.pointwise.data
3   @show PSIS_m5_3t.pointwise.statistic
4 end

```

```

PSIS_m5_3t.pointwise.data = 1:50
PSIS_m5_3t.pointwise.statistic = [:cv_elpd, :naive_lp
d, :p_eff, :mcse, :pareto_k]

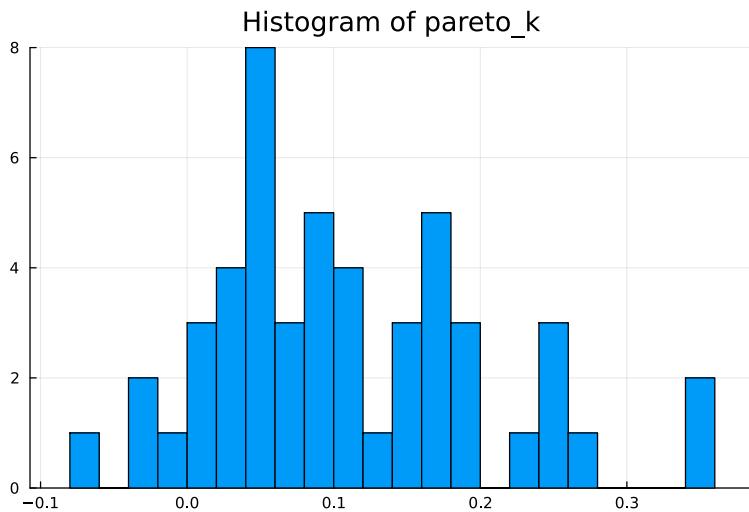
```

\emptyset				
(This table has no columns)				
1	:cv_elpd	-1.96683		
2	:naive_lpd	-1.89794		
3	:p_eff	0.0688969		
4	:mcse	0.0100315		
5	:pareto_k	0.0489091		

```

1 #2nd data point (divorce dataset)
2 PSIS_m5_3t.pointwise[data=2]

```



```
1 histogram(collect(PSIS_m5_3t.pointwise[statistic=5]),
  bins=20, title="Histogram of pareto_k")
```

ArgumentError: invalid index: ":pareto_k" of type String

Selection deleted

Stack trace

Here is what happened, the most recent locations are first:

1. `to_index(i::String) @ [indices.jl:300`
2. `to_index(A::AxisKeys.KeyedArray{Float64, 2,
 NamedDims.NamedDimsArray{(:data, :statistic), Float64,
 2, Matrix{Float64}}, Tuple{UnitRange{Int64},
 Vector{Symbol}}}, i::String) @ [indices.jl:277`
3. `_to_indices1(A::AxisKeys.KeyedArray{Float64, 2,
 NamedDims.NamedDimsArray{(:data, :statistic), Float64,
 2, Matrix{Float64}}, Tuple{UnitRange{Int64},
 Vector{Symbol}}}, inds::Tuple{Base.OneTo{Int64}},
 I1::String) @ [indices.jl:359`
4. `to_indices @ indices.jl:354`
5. `to_indices @ selectors.jl:212`
6. `to_indices @ indices.jl:344`
7. `getindex(::AxisKeys.KeyedArray{Float64, 2,
 NamedDims.NamedDimsArray{(:data, :statistic), Float64,
 2, Matrix{Float64}}, Tuple{UnitRange{Int64},
 Vector{Symbol}}}, ::Function, ::String)
 @ struct.jl:75`
8. `#getindex#26 @ names.jl:85`
9. `getindex @ names.jl:82`
10. `This cell: line 2
 1 # this approach does not work.
 2 PSIS_m5_3t.pointwise[statistic=:pareto_k"]`

```
1 # this approach does not work.
2 PSIS_m5_3t.pointwise[statistic=:pareto_k"]
```

\emptyset		
(This table has no columns)		
1	1	0.117485
2	2	0.0489091
3	3	0.128208
4	4	0.166018
5	5	0.18981
6	6	0.049939
7	7	0.166855
8	8	0.0443371
9	9	0.351656
10	10	0.0808025
⋮ more		

```
1 PSIS_m5_3t.pointwise(Symbol("pareto_k"))
```

► (50)

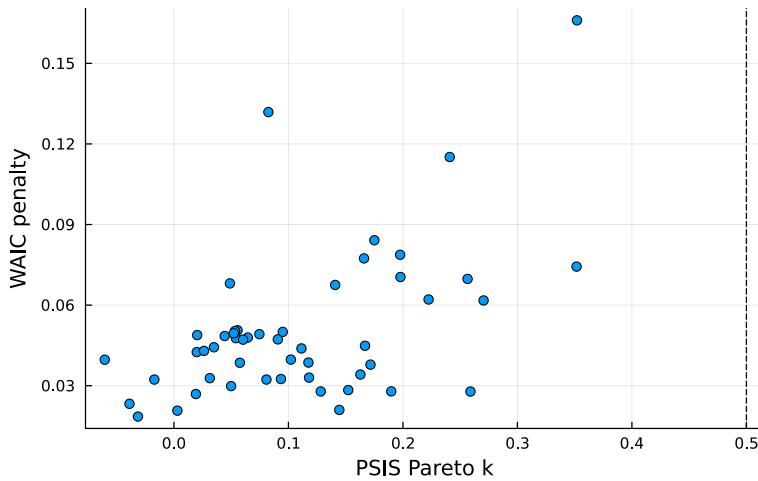
Selection deleted

```
1 begin
2     @show typeof(PSIS_m5_3t.pointwise(:cv_elpd))
3     size(PSIS_m5_3t.pointwise(:cv_elpd))
4 end
```

```
typeof(PSIS_m5_3t.pointwise(:cv_elpd)) = AxisKeys.K
eyedArray{Float64, 1, NamedDims.NamedDimsArray{(:data,), Float64, 1, SubArray{Float64, 1, Matrix{Float64}, Tuple{Base.Slice{Base.OneTo{Int64}}, Int64}, true}}, Base.RefValue{UnitRange{Int64}}}
```

- No outliers in PSIS K after replacing Normal() with student T

Student T (df=2) (m5.3)



```
1 begin
2     scatter(PSIS_m5_3t.pointwise(:pareto_k),
3             WAIC_m5_3t.penalty,
4             xlab="PSIS Pareto k", ylab="WAIC penalty",
5             title="Student T (df=2) (m5.3)")
6     vline!([0.5], c=:black, s=:dash)
7 end
```

	Location	Loc	Population	MedianAgeMarried
1	"District of Columbia"	"DC"	0.6	29.7
2	"Idaho"	"ID"	1.57	23.2
3	"Maine"	"ME"	1.33	26.4
4	"Utah"	"UT"	2.76	23.3

```
1 @time find_outliers(d_divorce,
PSIS_m5_3t.pointwise(:pareto_k), WAIC_m5_3t.penalty;
min_pareto_k=0.3, min_waic_penalty=0.1)
```

0.000195 seconds (162 allocations: 17.859 KiB)

7.36 m5_3t: 1D T-dist (failed in ForwardDiff but succeeded via using arraydist)

```
1 md"## 7.36 `m5_3t`: 1D T-dist (failed in ForwardDiff but
succeeded via using arraydist)"
```

TDist (generic function with 1 method)

```
1 TDist(μ, σ, ν) = μ + Distributions.TDist(ν)*σ
```

```

MethodError: no method matching
Float64(::ForwardDiff.Dual{ForwardDiff.Tag{DynamicPPL.DynamicPPLTag, Float64}, Float64, 4})
Closest candidates are:
(::Type{T})(::Real, !Matched)::RoundingMode) where T<:AbstractFloat
@ Base rounding.jl:207
(::Type{T})(::T) where T<:Number
@ Core boot.jl:792
Float64(!Matched)::Float32
@ Base float.jl:261
...

```

Stack trace

Here is what happened, the most recent locations are first:

1. `convert(::Type{Float64},`
`x::ForwardDiff.Dual{ForwardDiff.Tag{DynamicPPL.DynamicPPLTag, Float64}, Float64, 4}) @ [number.jl:7]`
- Selection deleted! `(A::Vector{Float64},`
`x::ForwardDiff.Dual{ForwardDiff.Tag{DynamicPPL.DynamicPPLTag, Float64}, Float64, 4}, i1::Int64)`
`@ [array.jl:1021]`
3. `macro expansion @ [This cell: line 10`
`9 for i in 1:no_of_D`
`10 μ[i] = a + bA * A[i] + bM * M[i]`
`11 D[i] ~ TDist(μ[i], σ, 2::Int64)`
4. `Show more...`

```

1 begin
2   @model function model_m5_3_singleT(A, M, D)
3     no_of_D = length(D)
4     σ ~ Exponential(1)
5     a ~ Normal(0, 0.2)
6     bA ~ Normal(0, 0.5)
7     bM ~ Normal(0, 0.5)
8     μ = similar(A)
9     for i in 1:no_of_D
10      μ[i] = a + bA * A[i] + bM * M[i]
11      D[i] ~ TDist(μ[i], σ, 2::Int64)
12    end
13  end
14
15  @time m5_3_singleT =
16    sample(model_m5_3_singleT(d_divorce.A, d_divorce.M,
17    d_divorce.D), NUTS(), 1_000);
18  m5_3_singleT_df = DataFrame(m5_3_singleT);
end

```

100%

	a	bA	bM	log_σ
1	0.00756524	-0.596769	-0.0244071	-0.405985
2	-0.016395	-0.734817	-0.0244059	-0.389975
3	0.0942605	-0.559865	0.0817123	-0.360903
4	-0.0454193	-0.894398	0.0496404	-0.36093
5	-0.0202127	-0.667483	-0.243533	-0.488368
6	-0.143304	-0.623603	0.36428	-0.459321
7	0.0517743	-0.744883	-0.170589	-0.384995
8	0.0122438	-0.751377	0.0110436	-0.62418
9	0.0583168	-0.550769	0.262727	-0.275354
10	-0.0616891	-0.739931	-0.169161	-0.352399
: more				
1000	-0.0218696	-0.894011	-0.041088	-0.525688

```

1 begin
2 #tdist_custom(μ) = TDist(μ, σ, 2::Int64)
Selection deleted
3 #lazyarray(f, x) = LazyArray(Base.broadcasted(f, x))
4 @model function model_m5_3_array_singleT(A, M, D)
5     #σ ~ Exponential(1)
6     log_σ ~ Normal(0, 0.3)
7     a ~ Normal(0, 0.2)
8     bA ~ Normal(0, 0.5)
9     bM ~ Normal(0, 0.5)
10    μ = @. a + bA * A + bM * M
11    D ~ arraydist(map(x -> TDist(x, exp(log_σ),
12        2::Int64), μ))
13 end
14
15 @time m5_3_array_singleT =
16 sample(model_m5_3_array_singleT(d_divorce.A,
17 d_divorce.M, d_divorce.D), NUTS(), 1_000);
18 m5_3_array_singleT_df = DataFrame(m5_3_array_singleT);
end

```

100%

① Found initial step size
 $\epsilon: 0.4$

6.361126 seconds (5.50 M allocations: 476.216 MiB, 2.85% gc time, 58.44% compilation time)

"stddev of σ is 0.08716211151314428"

1 est_summary(m5_3_array_singleT_df)

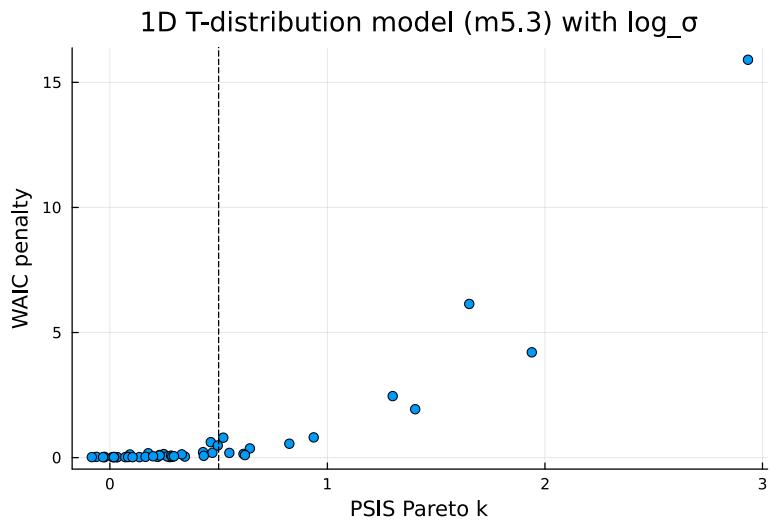
	exp(median(log_σ estimate)) 0.6382447917093729				
	o median est is 0.6382447935724113	mcmc_df_sum = 4x8 DataFrame	Row	variable	mean
x	nmissing	eltype	min	median	ma
oat64	Symbol	Float64	Float64	Float64	Fl
	Int64	DataType	Float64		
1	0.349162	0	0.0281614	-0.347638	0.027957
2	0.210936	0	-0.682299	-1.20396	-0.682835
3	0.797033	0	0.0479895	-0.551996	0.0432765
4	0.0270338	0	-0.450273	-0.879415	-0.449033

```
► (WAIC = [5.16654, 4.30301, 1.48751, 4.81252, 1.7165, 2.7939, 2.7
```

```
1 begin
2   @time m5_3_array_singleT_ll =
3     StatisticalRethinking.link(m5_3_array_singleT_df,
4       (r, (a,m,d)) -> StatsFuns.normlogpdf(r.a + r.bA * a
5         + r.bM * m, exp(r.log_σ), d),
6       zip(d_divorce.A, d_divorce.M, d_divorce.D));
7   m5_3_array_singleT_ll = hcat(m5_3_array_singleT_ll...);
8   @time PSIS_m5_3_array_singleT =
9     ParetoSmooth.psis_loo(ll_to_psis(m5_3_array_singleT_ll))
10  @time WAIC_m5_3_array_singleT =
11    psis.waic(m5_3_array_singleT_ll, pointwise=true)
12 end
```

- ⓘ No source provided for samples; variables are assumed to be from a Markov Chain. If the samples are independent, specify this with keyword argument 'source=:other'.

```
0.166821 seconds (1.06 M allocations: 25.030 MiB, 84.24% compilation time)
size(ll) = (1000, 50)
size(t) = (50, 1000)
0.010295 seconds (1.52 k allocations: 4.181 MiB)
0.000604 seconds (63 allocations: 790.562 KiB)
```



```
1 begin
2   scatter(PSIS_m5_3_array_singleT.pointwise(:pareto_k),
3     WAIC_m5_3_array_singleT.penalty,
4     xlabel="PSIS Pareto k", ylabel="WAIC penalty",
5     title="1D T-distribution model (m5.3) with log_σ")
6   vline!([0.5], c=:black, s=:dash)
7 end
```

	Location	Loc	Population	MedianAgeMarriage	MeanAgeDivorce
1	"Idaho"	"ID"	1.57	23.2	25
2	"Maine"	"ME"	1.33	26.4	13
3	"North Dakota"	"ND"	0.67	25.3	26
4	"Utah"	"UT"	2.76	23.3	29
5	"Wyoming"	"WY"	0.56	24.2	30

```
1 find_outliers(d_divorce,
PSIS_m5_3_array_singlet.pointwise(:pareto_k),
WAIC_m5_3_array_singlet.penalty; min_pareto_k=1.0,
min_waic_penalty=1.0)
```

- Many more outliers using 1D T-dist vs 2D T-dist.
- Primary cause: $\log\sigma$ estimates by 1D T-dist (-0.45) is much smaller than 2D T-dist (-0.15). But not sure why so different.

1 md"- Many more outliers using 1D T-dist vs 2D T-dist.
 2 - Primary cause: $\log\sigma$ estimates by 1D T-dist (-0.45) is
 much smaller than 2D T-dist (-0.15). But not sure why so
 different."

Try Truncated cauchy as σ (represented as $\log\sigma$ so that *estsummary* can be invoked).

```
1 md### Try Truncated cauchy as  $\sigma$  (represented as  $\log\sigma$  so  

  that est_summary can be invoked)."
```

	a	bA	bM	log_σ
1	0.053781	-0.841872	-0.118816	0.474882
2	-0.0709019	-0.726279	-0.0931306	0.553218
3	-0.120763	-0.605659	-0.00644269	0.571397
4	0.0110662	-0.904736	-0.214782	0.655921
5	-0.109545	-0.819336	-0.280735	0.573337
6	0.197469	-0.583952	0.462479	0.630095
7	0.224469	-0.41519	0.410673	0.727186
8	0.230619	-0.554135	0.0743648	0.651612
9	0.00885437	-0.73151	0.2222	0.48228
10	0.0831263	-0.732038	-0.185884	0.659455
: more				
1000	0.0279407	-0.566284	-0.0881914	0.60237

```

1 begin
2     #tdist_custom(μ) = TDist(μ, σ, 2::Int64)
3     #lazyarray(f, x) = LazyArray(Base.broadcasted(f, x))
4     @model function model_m5_3_array_singleT_cauchy(A, M, D)
5         #σ ~ Exponential(1)
6         log_σ ~ truncated(Cauchy(0, 1), 0, Inf)
7         a ~ Normal(0, 0.2)
8         bA ~ Normal(0, 0.5)
9         bM ~ Normal(0, 0.5)
10        μ = @. a + bA * A + bM * M
11        D ~ arraydist(map(x -> TDist(x, log_σ, 2::Int64),
12                           μ))
13    end
14
15    @time m5_3_array_singleT_cauchy =
16    sample(model_m5_3_array_singleT_cauchy(d_divorce.A,
17                                             d_divorce.M, d_divorce.D), NUTS(), 1_000);
18    m5_3_array_singleT_cauchy_df =
19    DataFrame(m5_3_array_singleT_cauchy);
end

```

100%

ⓘ Found initial step size
 $\epsilon: 0.8$

9.418990 seconds (8.01 M allocations: 659.759 MiB, 2.66% gc time, 71.78% compilation time)

```
"stddev of σ is 0.16799673935612133"
1 est_summary(m5_3_array_singleT_cauchy_df)
```

exp(median(log_σ estimate) 1.7785216766467562
σ median est is 1.778521743905237
mcmc_df_sum = 4×8 DataFrame
Row | variable mean min median ma
x | nmissing eltype σ
oat64 | Symbol Float64 Float64 Float64 Fl
Int64 | DataType Float64

1 | a 0.0260662 -0.272912 0.0262593
0.319545 | 0 Float64 0.0988683
2 | bA -0.694552 -1.23316 -0.69386 -
0.0608241 | 0 Float64 0.149136
3 | bM 0.0526198 -0.539655 0.0438684
0.817846 | 0 Float64 0.206826
4 | log_σ 0.583859 0.400195 0.575783
1.09432 | 0 Float64 0.0899607

- σ estimate (0.58) is very similar to the prior σ estimate (0.64).
- Probably as many outliers as before.

```
1 md"- σ estimate (0.58) is very similar to the prior σ
2 estimate (0.64).
- Probably as many outliers as before."
```

7.37 m5_1t: m5_1 but uses T-distribution

```
1 md"## 7.37 `m5_1t`: `m5_1` but uses T-distribution"
```

	a	bA	log_σ
1	0.0309265	-0.433028	0.226964
2	0.173826	-0.643376	-0.0378036
3	-0.238171	-0.670755	5.86626e-5
4	-0.109542	-0.67757	-0.162775
5	0.109961	-0.534454	-0.252549
6	-0.021251	-0.668724	0.251216
7	0.0391024	-0.52464	-0.143188
8	0.0684505	-0.596609	0.0976087
9	-0.0992144	-0.34197	-0.197242
10	-0.0550326	-0.495661	0.0724789
: more			
1000	0.0386741	-0.769246	0.0050955

```

1 begin
2   @model function model_m5_1t(A, D)
3     a ~ Normal(0, 0.2)
4     bA ~ Normal(0, 0.5)
5     log_σ ~ Normal(0, 0.2)
6     μ = @. a + bA * A
7     # The vector . below is optional for NUTS sample.
8     # But psis_loo() requires it.
9     D ~ IsoTDist(2, μ, exp(log_σ))
10    end
11
12  @time m5_1t = sample(model_m5_1t(d_divorce.A,
13    d_divorce.D), NUTS(), 1000)
14  m5_1t_df = DataFrame(m5_1t);
15 end

```

100%

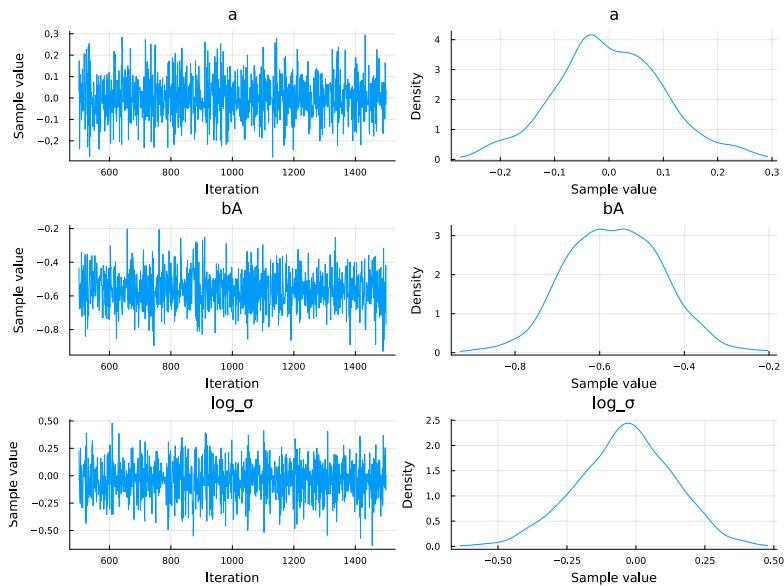
① Found initial step size
ε: 0.05

4.209546 seconds (3.30 M allocations: 265.811 Mi B, 1.46% gc time, 70.28% compilation time)

"stddev of σ is 0.16707628322371754"

1 est_summary(m5_1t_df; σ_ind=3)

	exp(median(log_σ estimate)) 0.965735554601083			
	σ median est is 0.9657355548513866	mcmc_df_sum = 3x8 DataFrame		
Row	variable	mean	min	median
max	nmissing	eltype	σ	Float64
	Symbol	Float64	Float64	Float64
Float64	Int64	DataType	Float64	
1	a	0.00118648	-0.27642	-0.00260447
0.294911		0	Float64	0.0995105
2	bA	-0.565421	-0.929385	-0.564129
-0.202886		0	Float64	0.11323
3	log_σ	-0.0392006	-0.637995	-0.0348652
0.479218		0	Float64	0.172446



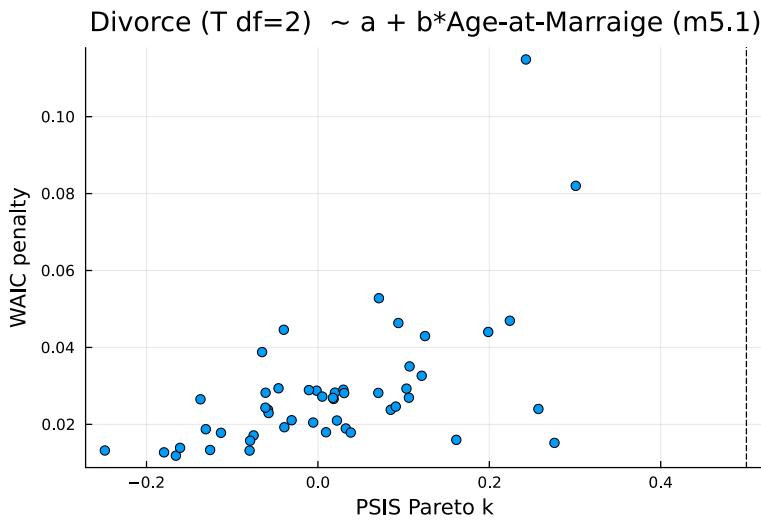
```
1 plot(m5_1t)
```

```
► (WAIC = [4.04416, 3.70151, 2.42175, 4.03876, 2.57078, 3.12294, 3
```

```
1 begin
2   m5_1t_ll = StatisticalRethinking.link(m5_1t_df,
3     (r, (a,d)) -> Distributions.logpdf(IsoTDist(2, [r.a
4       + r.bA * a], exp(r.log_σ)), [d]),
5     zip(d_divorce.A, d_divorce.D));
6   m5_1t_ll = hcat(m5_1t_ll...);
7
8   m5_1t_t = ll_to_psis(m5_1t_ll)
9   @show size(m5_1t_t)
10  PSIS_m5_1t = ParetoSmooth.psis_loo(m5_1t_t)
11  WAIC_m5_1t = psis.waic(m5_1t_ll, pointwise=true)
12 end
```

- ⓘ No source provided for samples; variables are assumed to be from a Markov Chain. If the samples are independent, specify this with keyword argument 'source=:other'.

```
size(ll) = (1000, 50)
size(t) = (50, 1000)
size(m5_1t_t) = (50, 1000, 1)
```



```

1 begin
2     scatter(PSIS_m5_1t.pointwise(:pareto_k),
3             WAIC_m5_1t.penalty,
4             xlab="PSIS Pareto k", ylab="WAIC penalty",
5             title="Divorce (T df=2) ~ a + b*Age-at-Marraige
(m5.1)")
6             vline!([0.5], c=:black, s=:dash)
7 end

```

- If σ is modelled by $\text{Exponential}(1)$, stddev of σ of `m5_1t_t` is much larger than `m5_1` (Normal distribution), which may explain why there are `more Pareto K outliers` in this student T model than Normal.
- But once $\log_{\sigma} \sim \text{Normal}(0, 0.2)$ is used instead, T distribution is better.
- A similar model by PyMC (another notebook) produces a much smaller stddev, thus no Pareto K outlier. Not sure why.

```

1 md"- If  $\sigma$  is modelled by  $\text{Exponential}(1)$ , stddev of  $\sigma$  of
'm5_1t_t' is much larger than 'm5_1' (Normal distribution),
which may explain why there are 'more Pareto K outliers' in
2 this student T model than Normal.
3 - But once  $\log_{\sigma} \sim \text{Normal}(0, 0.2)$  is used instead, T
distribution is better.
- A similar model by PyMC (another notebook) produces a
much smaller stddev, thus no Pareto K outlier. Not sure
why."

```

	Location	Loc	Population	MedianAgeMarriage	MeanAgeDivorce
1	"Arkansas"	"AR"	2.92	24.3	26
2	"Idaho"	"ID"	1.57	23.2	25
3	"Minnesota"	"MN"	5.3	26.3	15
4	"North Dakota"	"ND"	0.67	25.3	26
5	"Rhode Island"	"RI"	1.05	28.2	15
6	"Utah"	"UT"	2.76	23.3	29

```
1 @time find_outliers(d_divorce,
PSIS_m5_1t.pointwise(:pareto_k), WAIC_m5_1t.penalty;
min_pareto_k=0.22, min_waic_penalty=0.05)
```

0.000205 seconds (162 allocations: 18.516 KiB)